

Concurrency and Parallelism

Roadmap

- In the following sections,
- we'll first turn our attention to implicit parallelism, and how best to optimize our software to take advantage of it.
- Next, we'll review the most common forms of explicit parallelism. Then we'll explore the various concurrent programming techniques used to harness explicitly parallel computing platforms.
- Finally, we'll round out the discussion of parallel programming by discussing SIMD vector processing, and how it applies to GPU design and general-purpose GPU programming (GPGPU) techniques.

The Advent of Parallelism and Concurrency

- 1970s
 - Intel 8087 FP coprocessor: 50 kFLOPS (5.0×10^4)
 - Cray-1 supercomputer: 160 MFLOPS (1.6×10^8)
- Today
 - AMD Jaguar x86-64 8-core: 100 GFLOPS (1.0×10^{11})
 - China's Sunway TaihuLight: 93 PFLOPS (9.3×10^{16})
 - 7 to 8 orders of magnitude improvement!
 - Currently, **Summit is the fastest supercomputer**
 - **Speed:** 148.6 petaFLOPS
 - **Cores:** 2,414,592
 - **Vendor:** IBM
 - **Location:** Oak Ridge National Laboratory, United States

Overview

- Today's computer hardware is increasingly **parallel**
- Need to understand **concurrency** to take advantage of parallelism
- Need to understand the **hardware** and the **kernel** to understand concurrent programming!
- In this talk we'll cover:
 - modern **parallel** computing **hardware**,
 - **concurrent programming** techniques, and
 - how to apply this to building a **high-performance game engine**

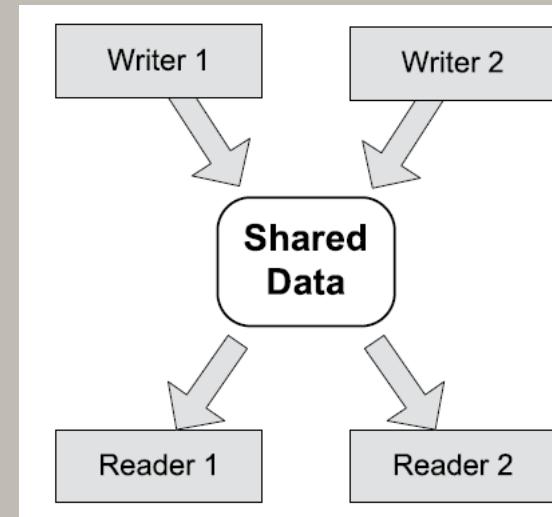
What is Concurrency?

- What is concurrency?
- What is parallelism?
- How are they related? How do they differ?

Discussion: 5 mins

What is Concurrency?

- Concurrency can be defined as the **transformation** of a *shared data file*
 - Shared by multiple **flows of control**
 - processes,
 - threads,
 - fibers...
 - Multiple **readers** and/or **writers**
 - ... of *shared* data!
- Did I mention the data has to be **shared**?
- *race condition (data race)*
- two or more flows of control “compete” to be the first to read, modify and write a chunk of shared data

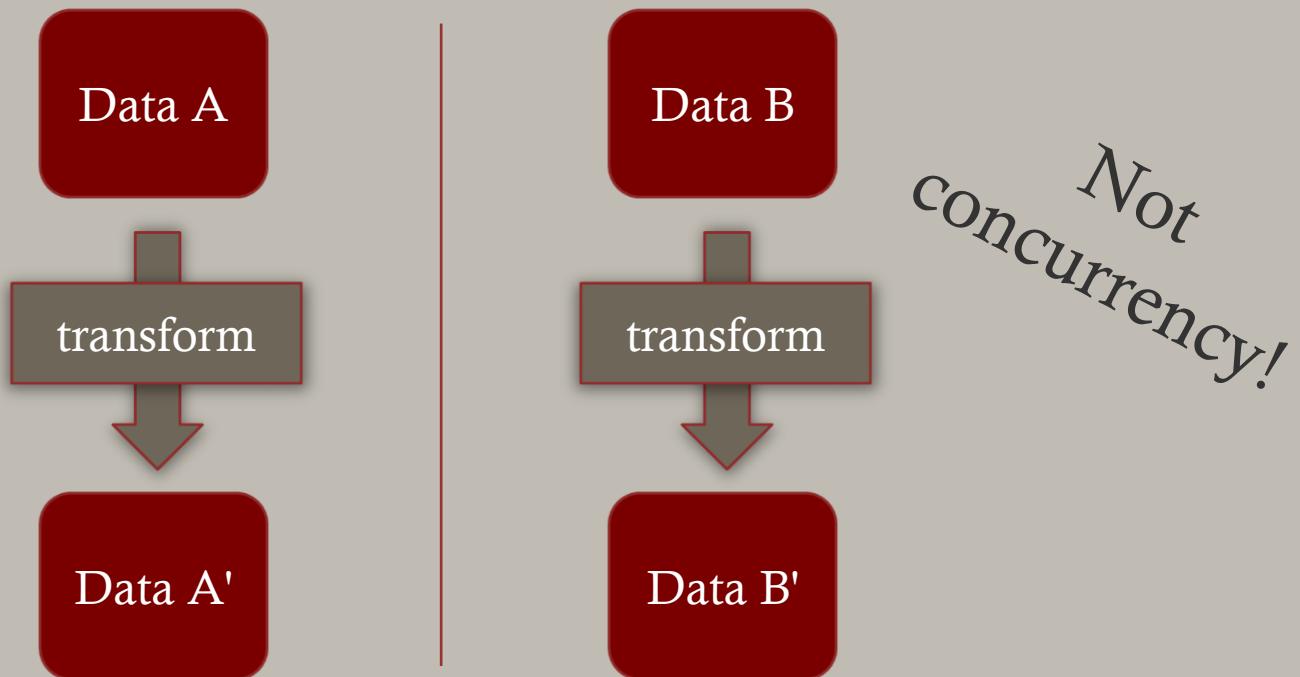


What is Concurrency?

- Data “file” can **be** anything
 - A **global Boolean variable** shared between threads
 - A shared **queue**
- Data file can **live** anywhere
 - Virtual memory space within a **multithreaded process**
 - Multiple processes **sharing virtual memory pages** within a single computer
 - GPU and CPU(s) **sharing physical RAM**
 - e.g., PS4, Xbox One
 - A **pipe** between two processes (`cat foo.txt | grep "bar"`)
 - A file living on a **network drive** that’s accessible by multiple computers

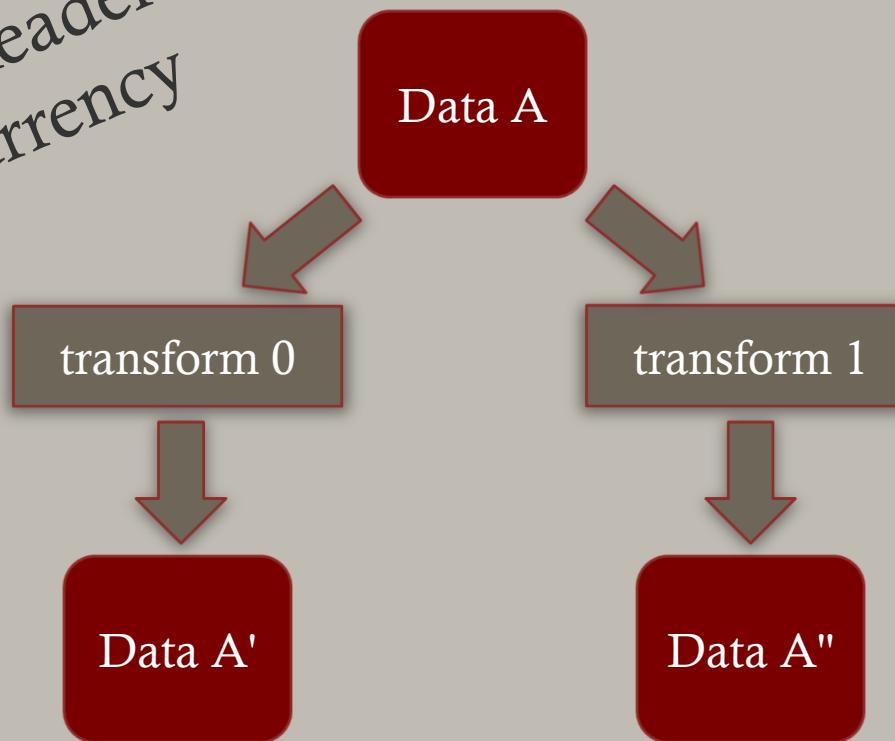
What is Concurrency?

- If data **isn't being shared**, it **isn't concurrency**
 - It's just computing “at the same time”

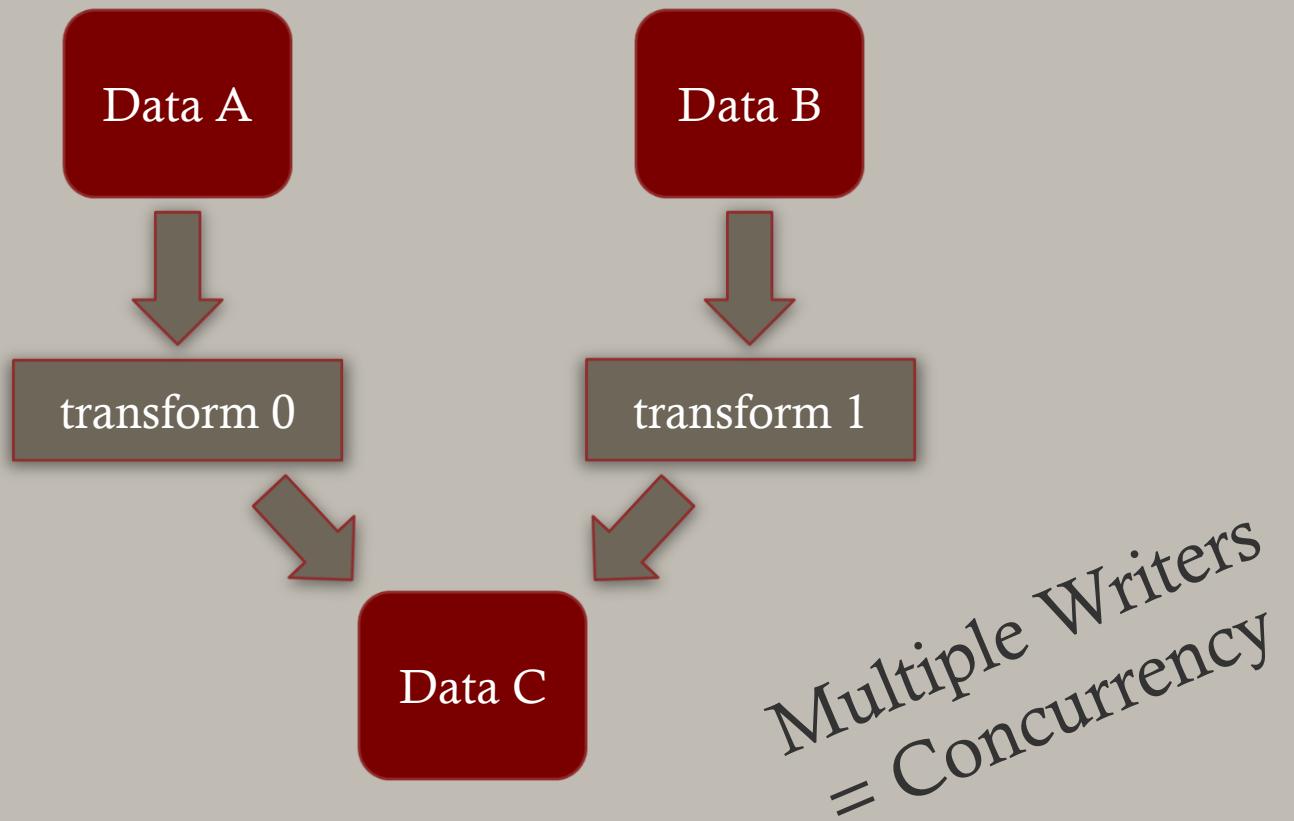


What is Concurrency?

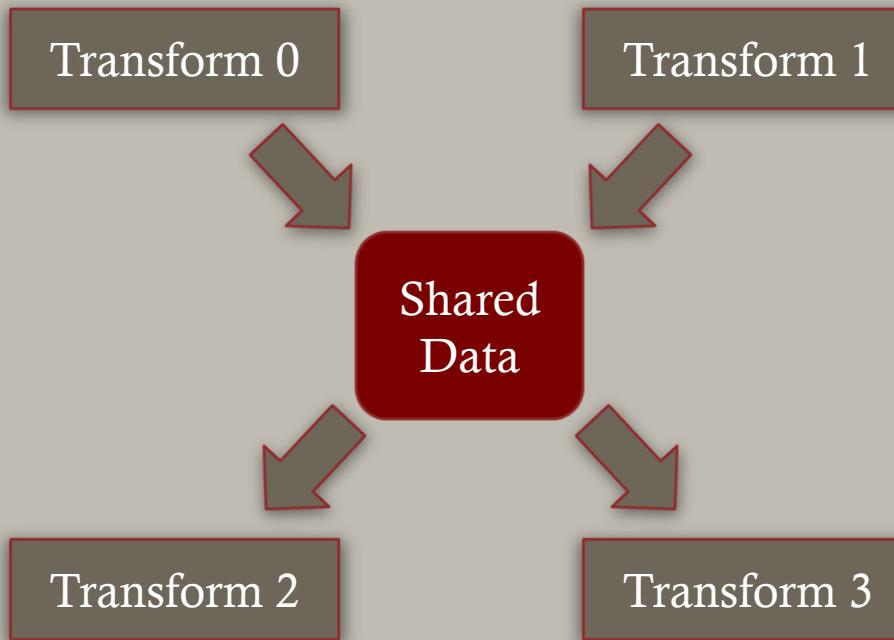
Multiple Readers
= Concurrency



What is Concurrency?



Multiple Writers
= Concurrency



Multiple Readers
= Concurrency

Concurrency versus Parallelism

- **Concurrency** means multiple flows of control operating on shared data
- **Parallelism** is multiple hardware components operating simultaneously
 - *Concurrency* is possible even on a **single-core** CPU
 - e.g., **preemptive** multitasking
 - Likewise, *parallel hardware* can improve performance of even **single-threaded** code
 - implicit parallelism (**pipelined** or **superscalar** CPU architectures)

The Advent of Parallelism and Concurrency

- In 1965, Intel co-founder Gordon Moore noticed:
 - computing power, measured in **transistors per unit area** of silicon, had been roughly **doubling** every year since the invention of the digital computer
- Moore’s “Law” states that:
 - At a given price point...
 - amount of computing power **doubles** roughly every **18 months**
- Reality is a bit more complex...

History of Processor Performance

Source: https://booksite.elsevier.com/9780123838728/figures/PDF/Chapter_01.pdf

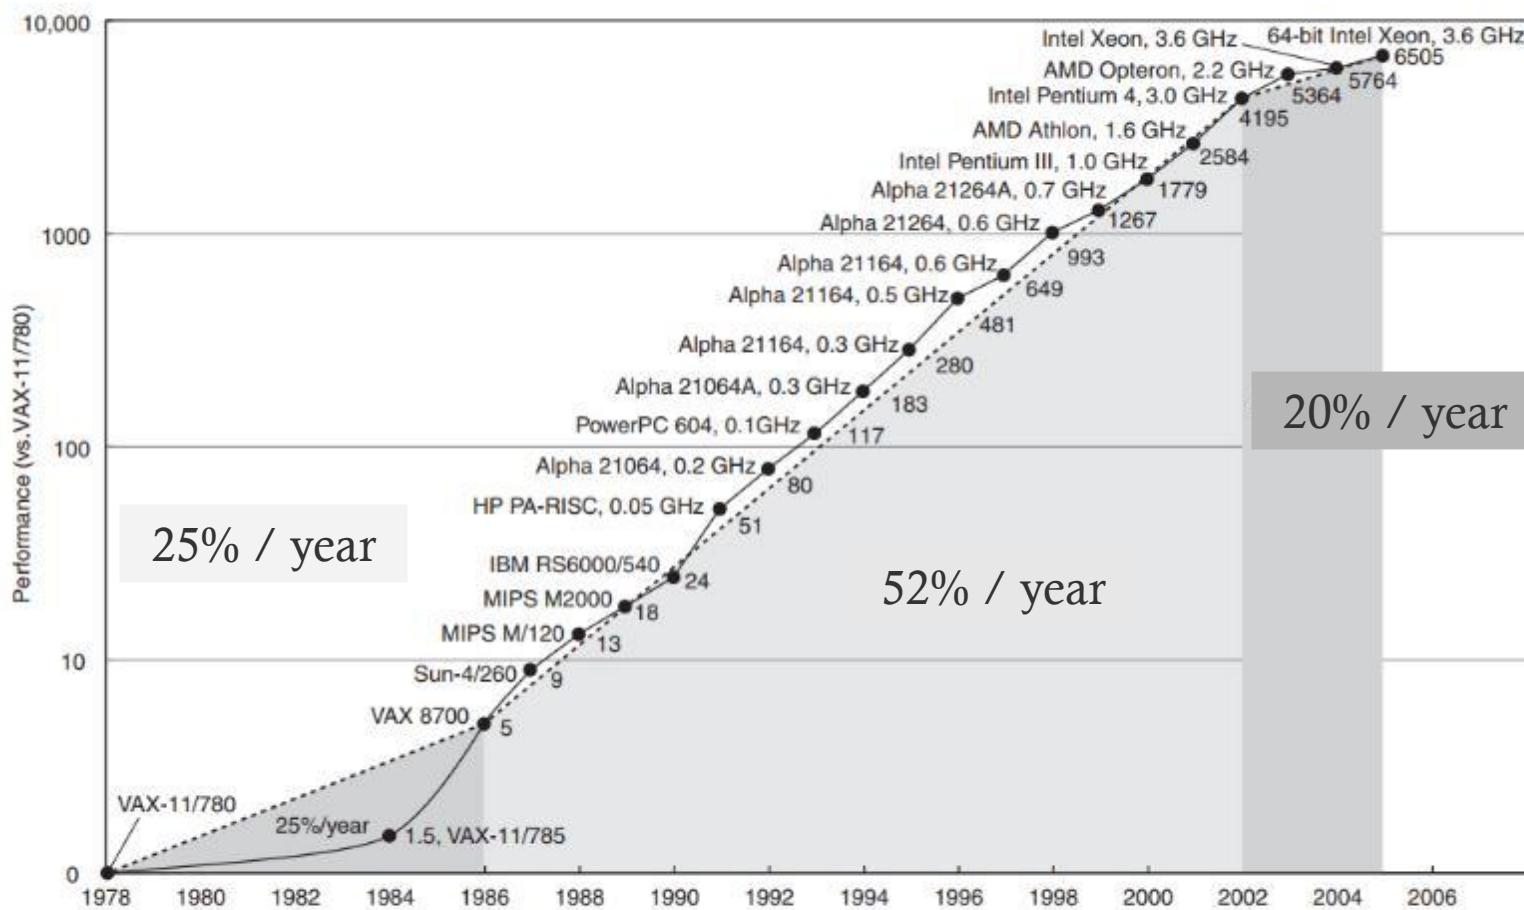


FIGURE 1.16 Growth in processor performance since the mid-1980s. This chart plots performance relative to the VAX 11/780 as measured by the SPECint benchmarks (see Section 1.8). Prior to the mid-1980s, processor performance growth was largely technology-driven and averaged about 25% per year. The increase in growth to about 52% since then is attributable to more advanced architectural and organizational ideas. By 2002, this growth led to a difference in performance of about a factor of seven. Performance for floating-point-oriented calculations has increased even faster. Since 2002, the limits of power, available instruction-level parallelism, and long memory latency have slowed uniprocessor performance recently, to about 20% per year. Copyright © 2009 Elsevier, Inc. All rights reserved.



History of Processor Performance

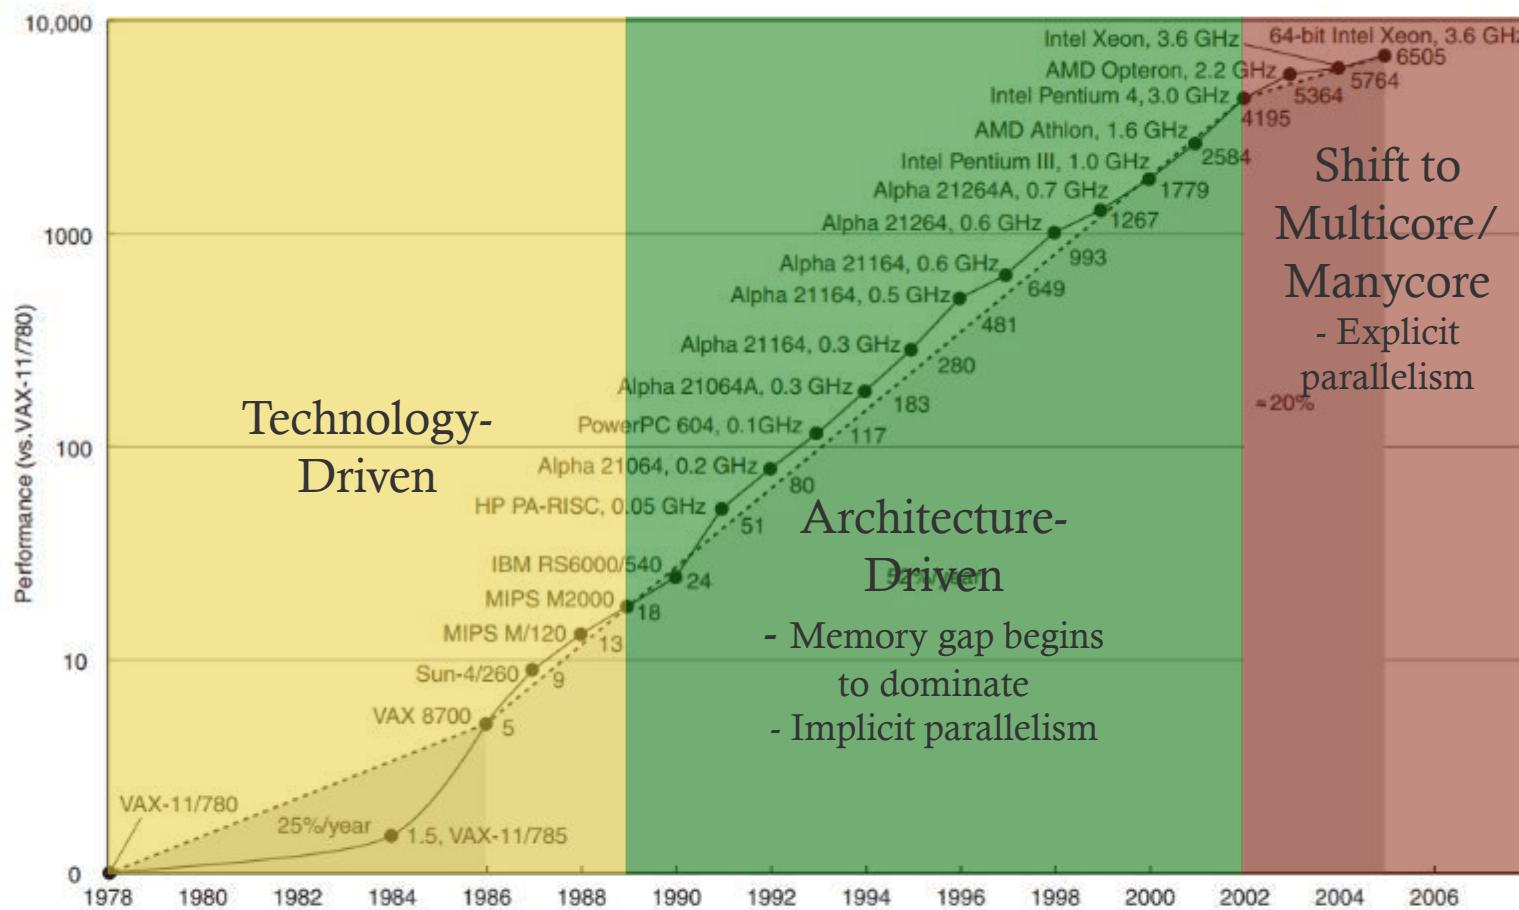
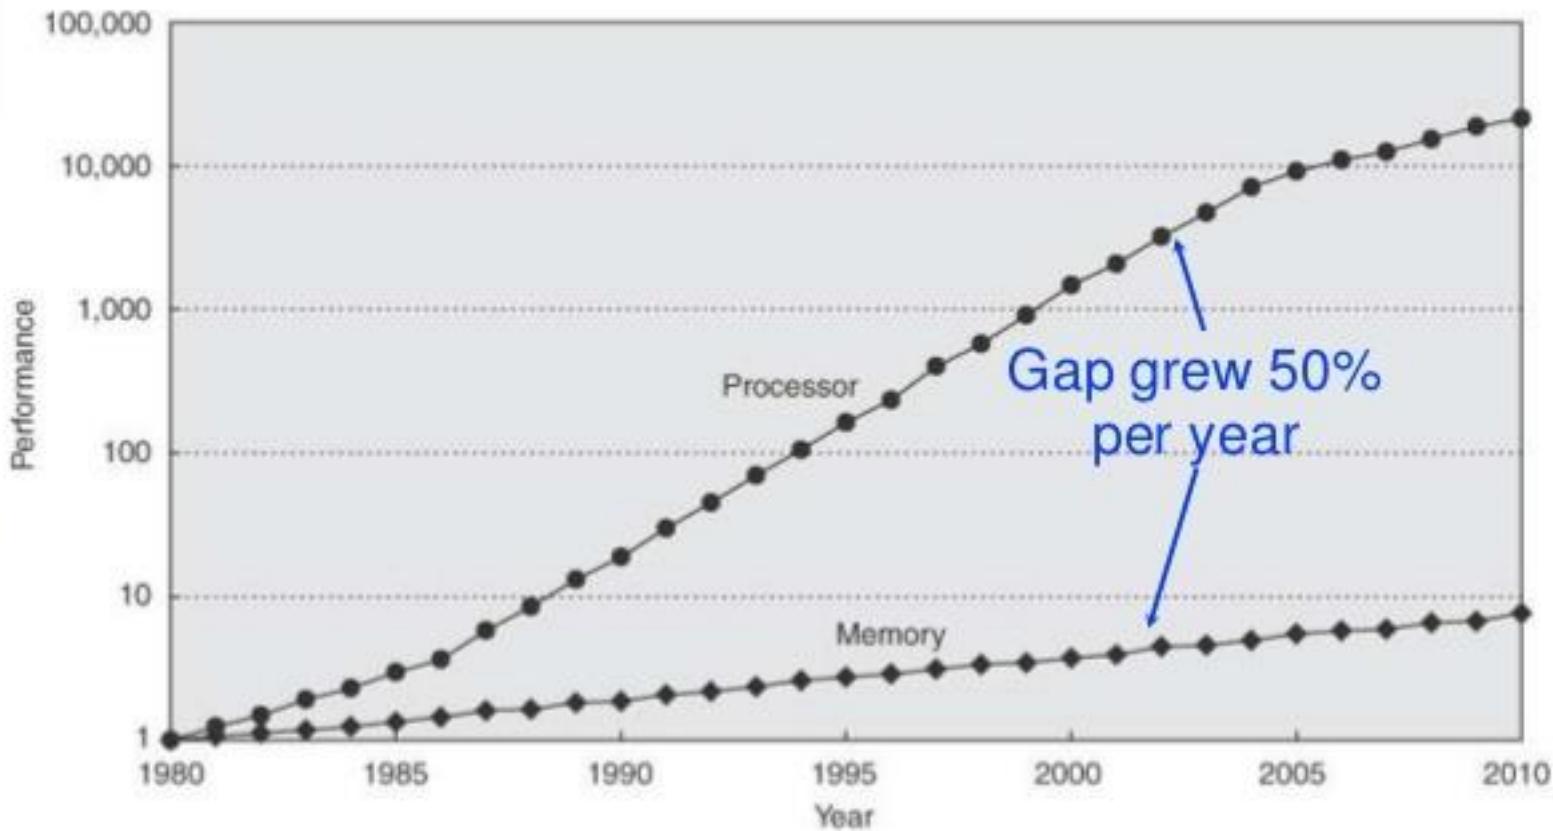


FIGURE 1.16 Growth in processor performance since the mid-1980s. This chart plots performance relative to the VAX 11/780 as measured by the SPECint benchmarks (see Section 1.8). Prior to the mid-1980s, processor performance growth was largely technology-driven and averaged about 25% per year. The increase in growth to about 52% since then is attributable to more advanced architectural and organizational ideas. By 2002, this growth led to a difference in performance of about a factor of seven. Performance for floating-point-oriented calculations has increased even faster. Since 2002, the limits of power, available instruction-level parallelism, and long memory latency have slowed uniprocessor performance recently, to about 20% per year. Copyright © 2009 Elsevier, Inc. All rights reserved.



Processor Memory Gap



© 2007 Elsevier, Inc. All rights reserved.

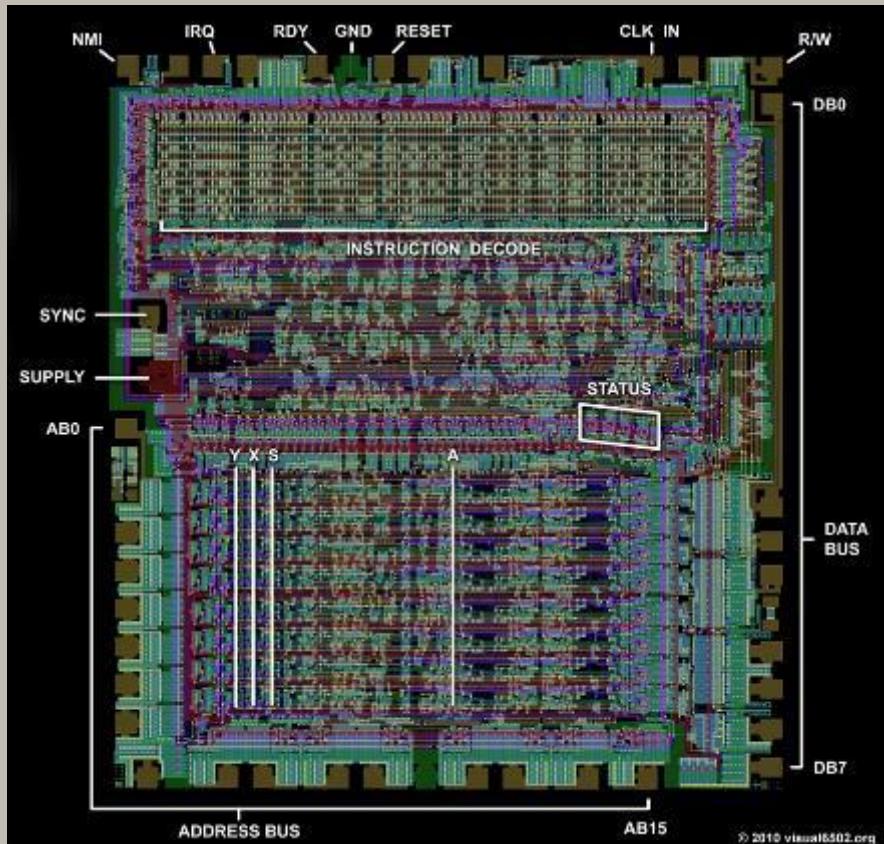
The Advent of Parallelism and Concurrency

- Increases in **clock speed** led to:
 - increased heat generation
 - increased power consumption
 - decreased signal-to-noise ratio
- Increase in **memory gap** led to:
 - Less bang for buck when improving single CPU core
- Increases in **transistor count** led to question:
 - How best to use all those transistors?

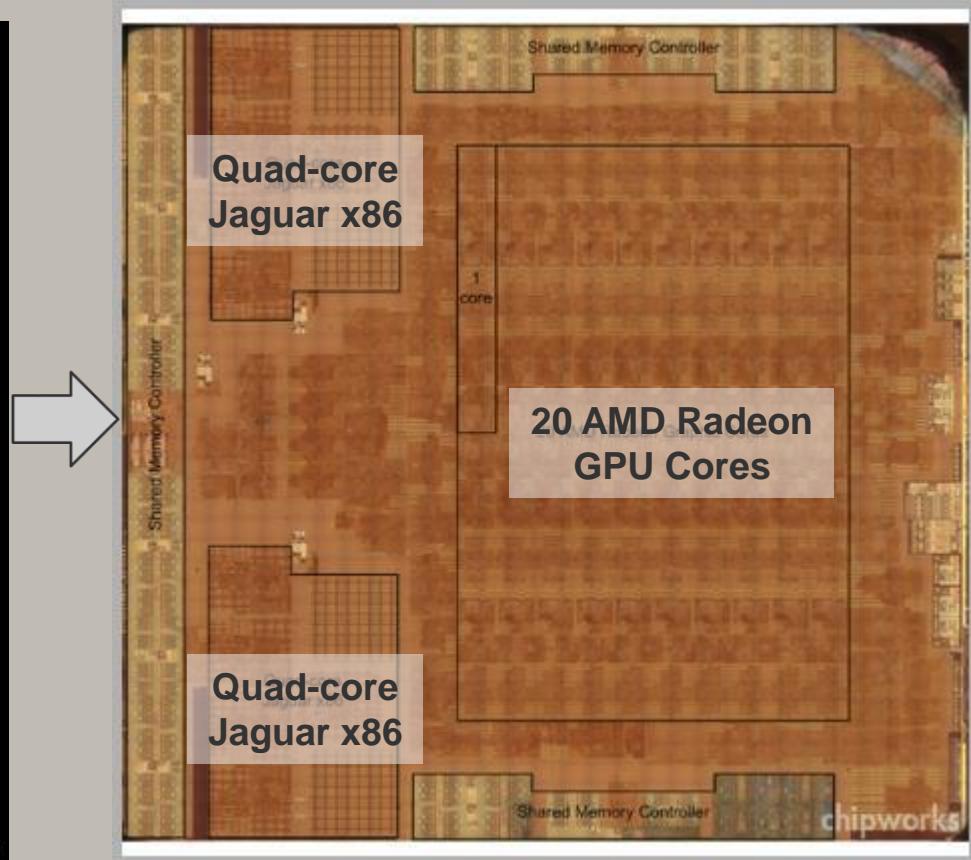
The Advent of Parallelism and Concurrency

- Result: Increased use of **parallelism** in consumer electronics

Single-core CMOS Tech 6502 (Apple II)



8-core AMD Jaguar + Radeon HD 7870 GPU (PS4)



The Advent of Parallelism and Concurrency

- Principles of parallelism known and in use since the **1960s**
- But back then, parallelism was reserved for high-end **supercomputers**
 - 1964: Seymour Cray pioneered **pipelining** with CDC 6600
 - 1972: CDC 8600 pioneered **multiprocessor** design by strapping four CDC 7600s together, executing in lock step
 - 1975: Cray-1 pioneered **vector processing** (SIMD)

Implicit Parallelism

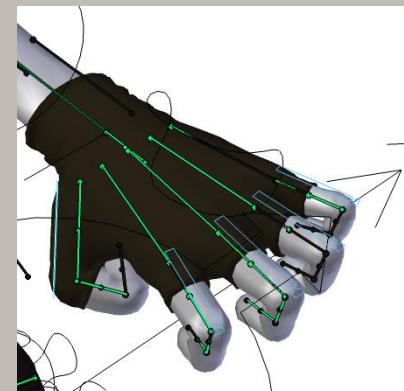
- **Implicit parallelism** is the use of parallelism to improve performance of **preexisting (sequential)** code
 - Little or no programmer/compiler intervention required
 - *although knowledge of implicit parallelism aids optimization!*
- Used in consumer electronics since ~1989
 - **Pipelining**
 - **Superscalar**
 - **VLIW**
 - e.g., pipelined CPU can run existing code faster than a non-pipelined CPU that supports the same ISA

Explicit Parallelism

- **Explicit parallelism** exposes the presence of parallel computing hardware to the:
 - programmer and/or
 - compiler
- Used in consumer electronics since ~2002
 - **Multiprocessor** computers
 - **Multicore** CPUs
 - SIMD **vector processing unit** in x86 (SSE)
 - **Manycore** GPUs (SIMT)
 - computer clusters
 - grid computing
 - cloud computing
- Explicit parallelism supports **concurrency**

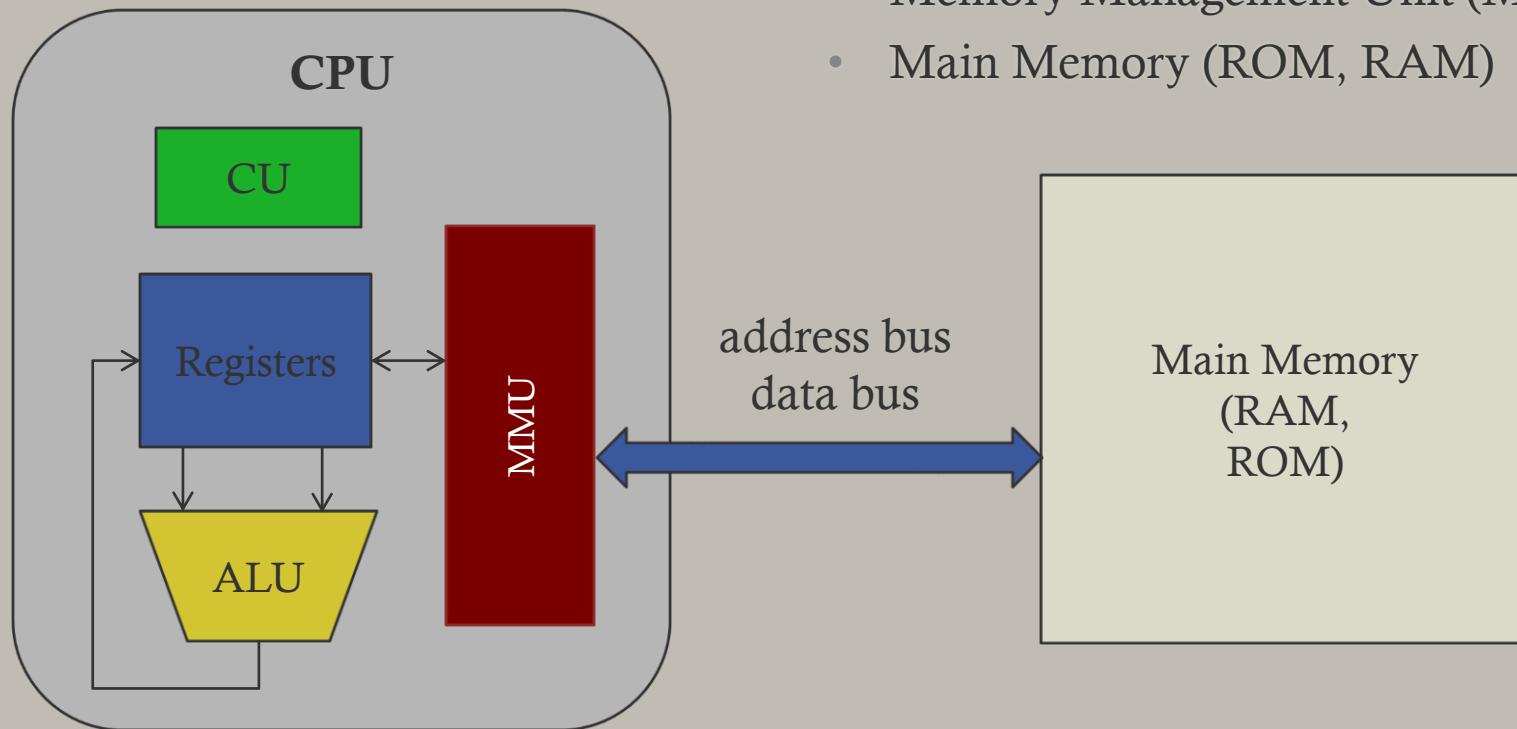
Task versus Data Parallelism

- *Task parallelism*: When multiple heterogeneous operations are performed in parallel.
 - Example: Performing animation calculations on one core while performing collision checks on another.
- *Data parallelism*: When a *single* operation is performed on multiple data items in parallel.
 - Example: Calculating 1000 skinning matrices by running 250 matrix calculations on each of four cores.



A Simple Computer

- Simplest von Neumann architecture
 - Control Unit (CU)
 - Arithmetic/Logic Unit (ALU)
 - Registers
 - Memory Management Unit (MMU)
 - Main Memory (ROM, RAM)



Flynn's Taxonomy

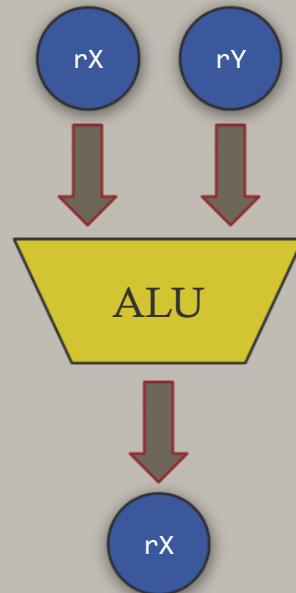
- Proposed by Michael J. Flynn of Stanford University in 1966
- Breaks parallelism into **instruction-level parallelism** (ILP) and **data-level parallelism** (DLP)
 - Single instruction, single data (SISD)
 - Multiple instruction, multiple data (MIMD)
 - Single instruction, multiple data (SIMD)
 - Multiple instruction, single data (MISD)
- Single instruction, multiple thread (SIMT): GPUs

Flynn's Taxonomy

SISD

Instruction Stream

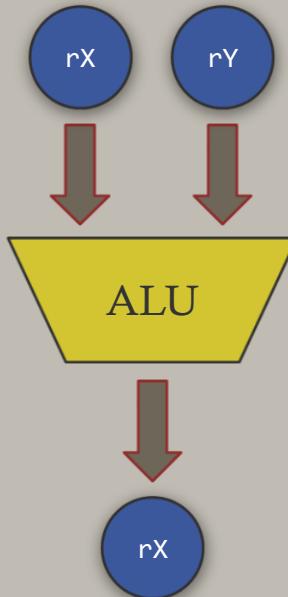
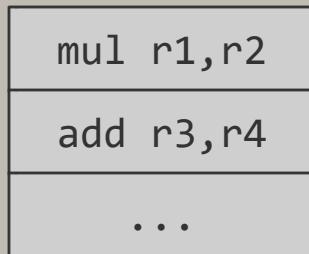
mul r1,r2
add r3,r4
...



Flynn's Taxonomy

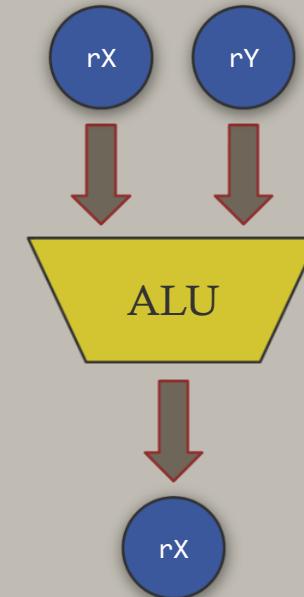
MIMD

Instruction Stream



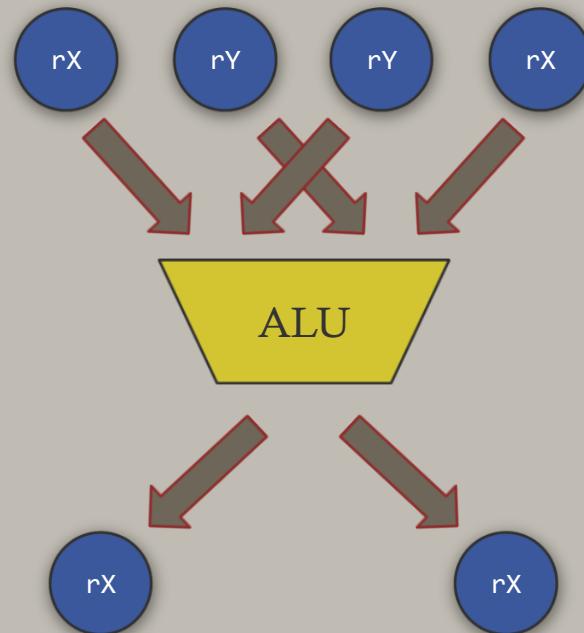
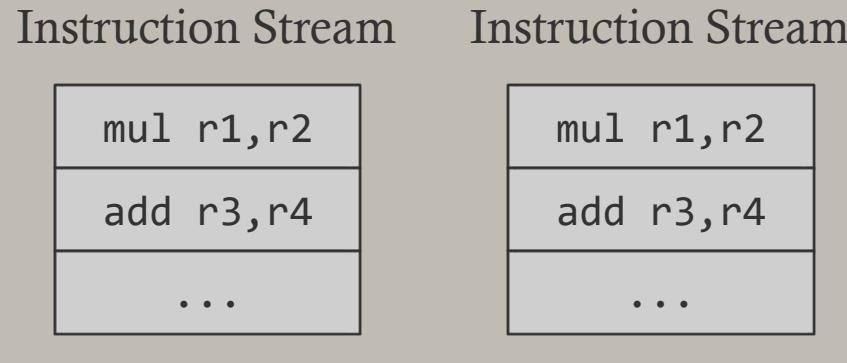
Instruction Stream

div r5,r2
sub r3,r1
...



Flynn's Taxonomy

MIMD
(time-sliced)



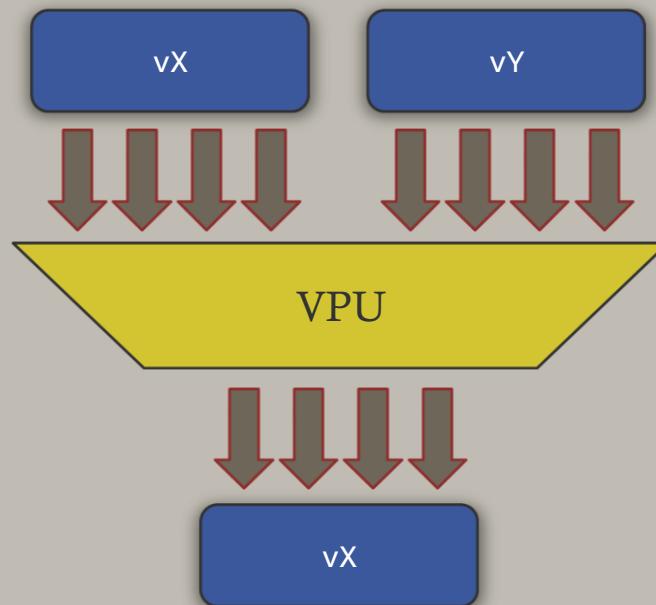
- Not parallelism...
- but this *is* concurrency

Flynn's Taxonomy

SIMD

Instruction Stream

mul v1,v2
add v3,v4
...

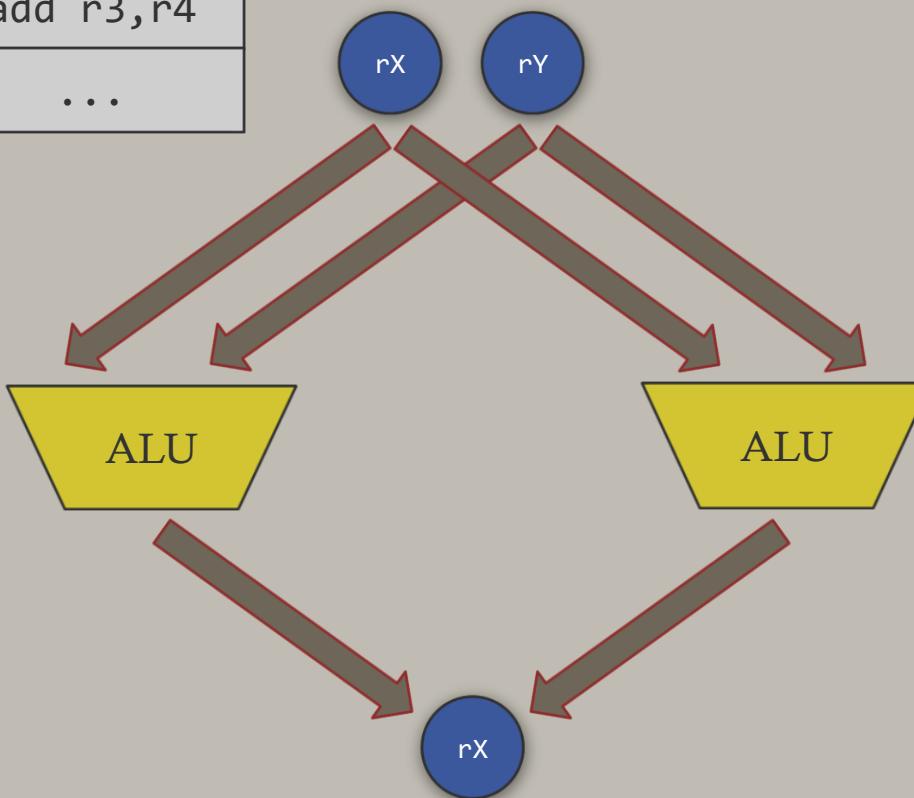


Flynn's Taxonomy

MISD

Instruction Stream

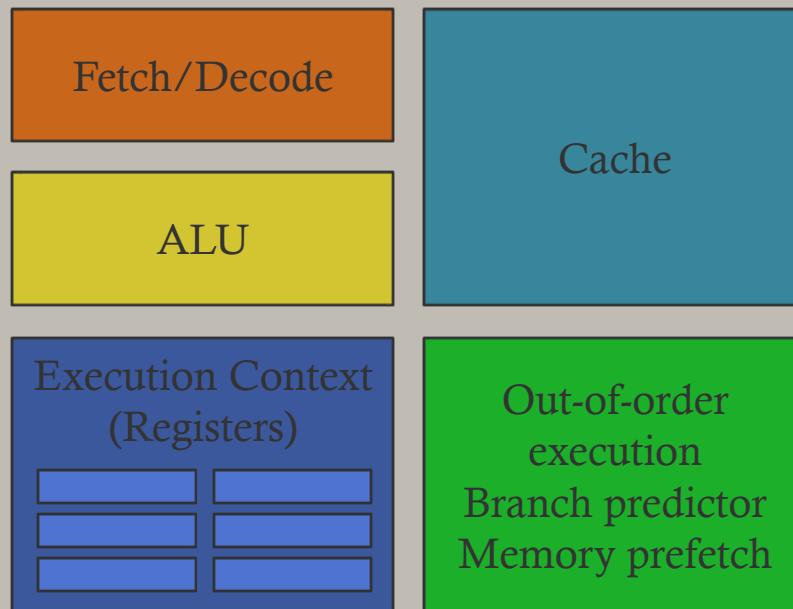
mul r1,r2
add r3,r4
...



- Primarily used for **fault tolerance** via redundancy

Manycore GPU (SIMT)

- Typical CPU core



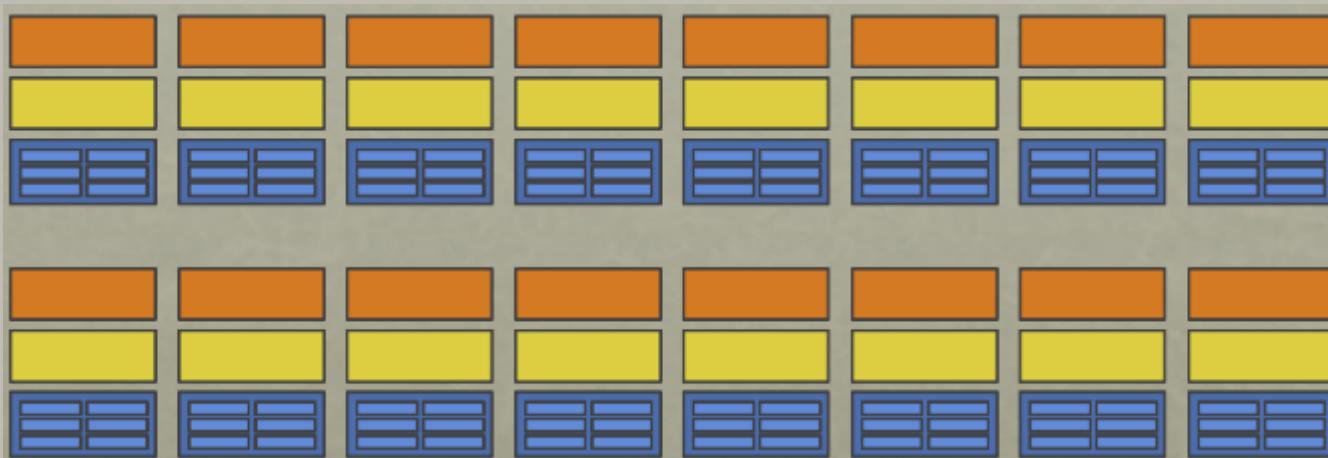
Manycore GPU (SIMT)

- Remove fancy OOO logic, branch predictor, cache



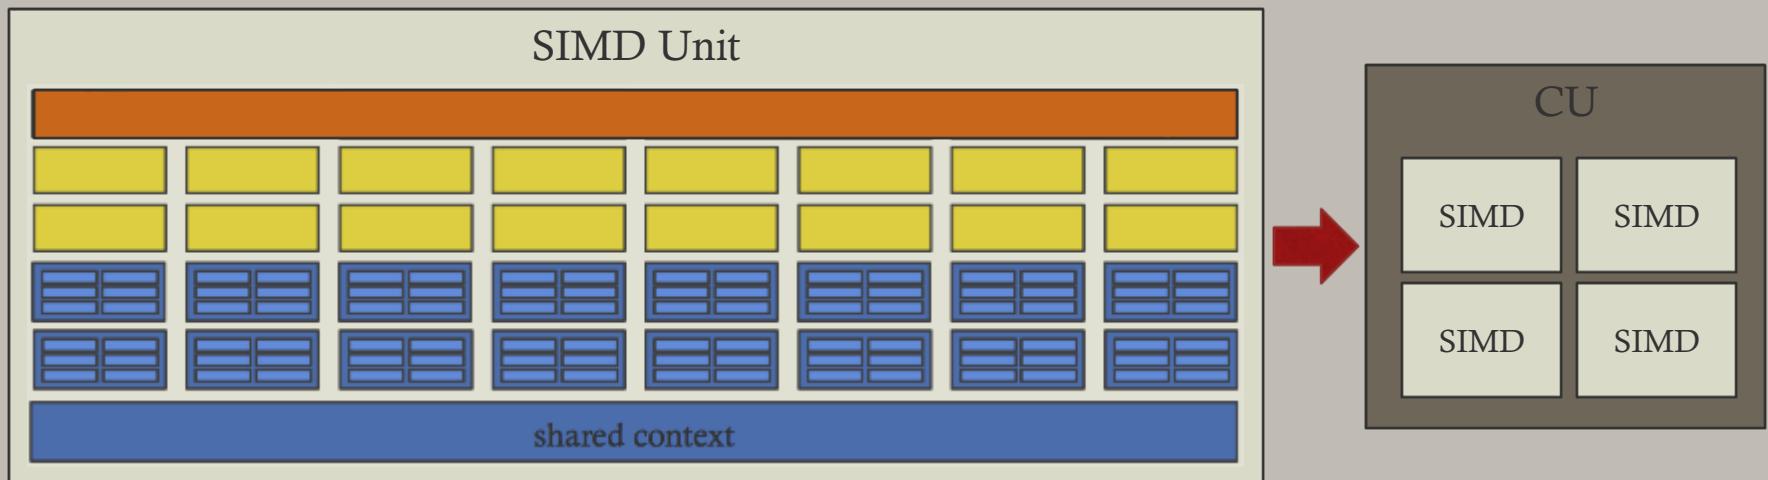
Manycore GPU (SIMT)

- Duplicate 16 times



Manycore GPU (SIMT)

- Group 16 ALUs and 16 execution contexts, managed by one fetch/decode unit to make a SIMD unit
- Four SIMDs per compute unit (CU) = 64 “lanes”



IMPLICIT PARALLELISM

Pipelining

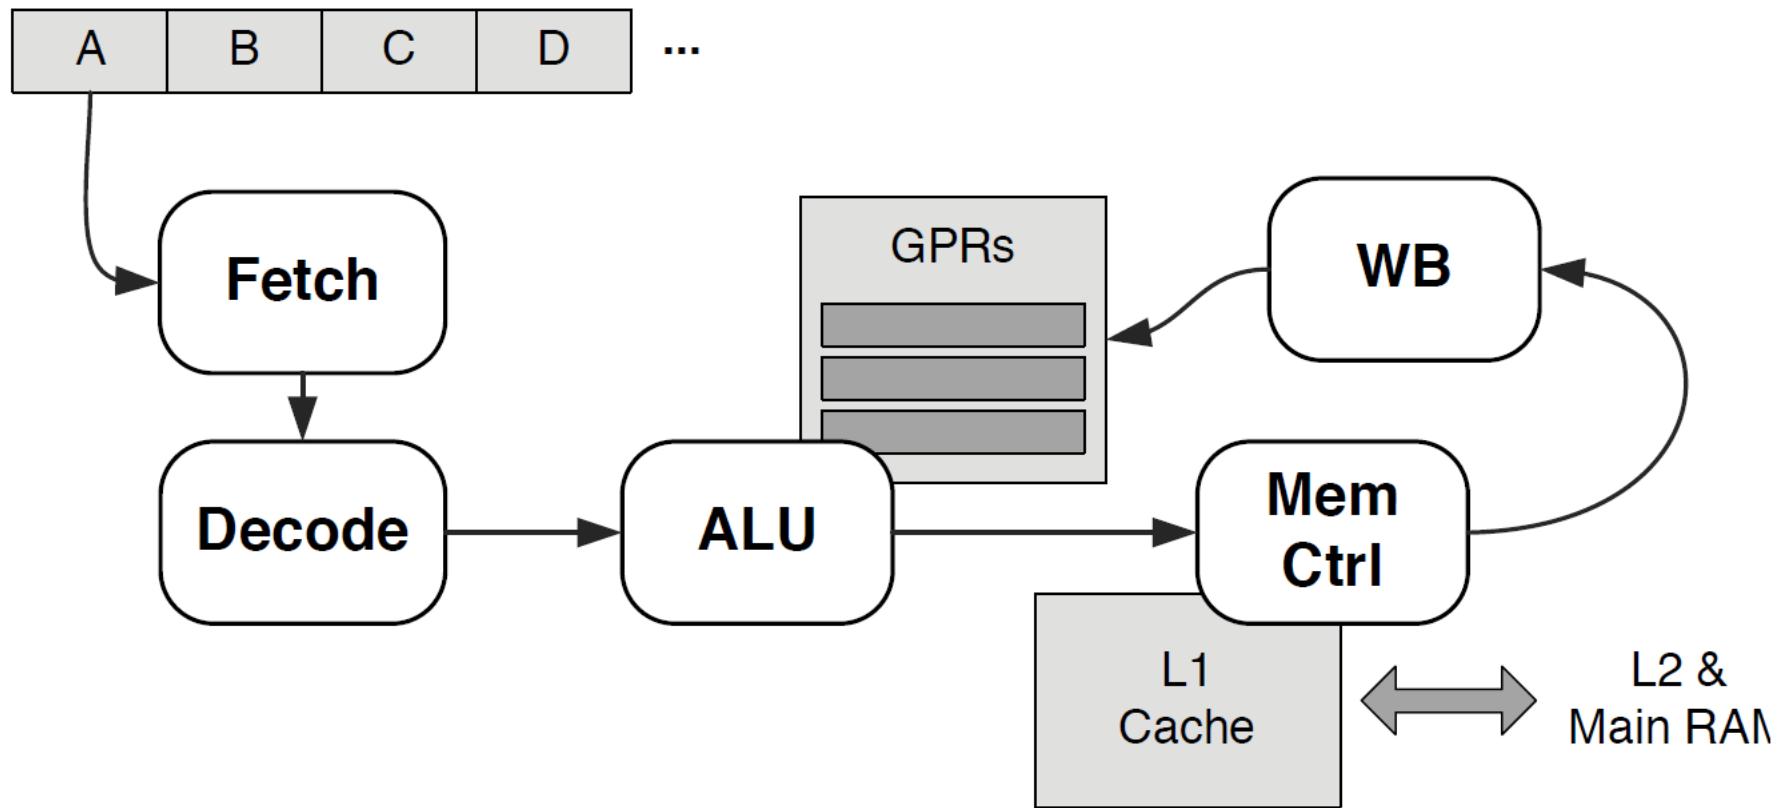
- Early CPUs executed **serially**
 - One instruction at a time
 - Components within CPU not well separated
 - Many components **idle** while an instruction executed
- Modern CPUs are **pipelined**
 - Each instruction passes through clearly-defined **stages**, each corresponding to a component within CPU core
 - Clear separation between components
 - Can keep all components busy by allowing multiple instructions to be “in flight” simultaneously
 - Each using a different component within the core

Pipelining

- Different CPUs define their stages differently
 - From four to five stages...
 - ... to up to 30+ stages in deeply pipelined CPUs
- Basic stages:
 1. **Fetch** an instruction word from memory (**F**)
 2. **Decode** instruction into opcode and operands (**D**)
 3. **Register access** (**R**)
 - We'll assume D and R are *combined* (**D/R**)
 4. **Execute** instruction on ALU (**E**)
 5. **Memory** read/write (**M**)
 6. **Register writeback**, i.e., store results in registers (**W**)

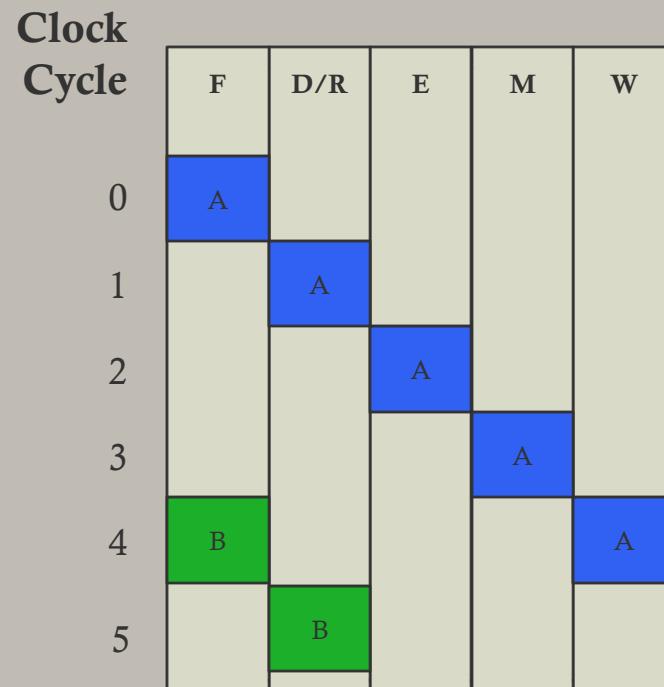
Pipelining

Instruction Stream



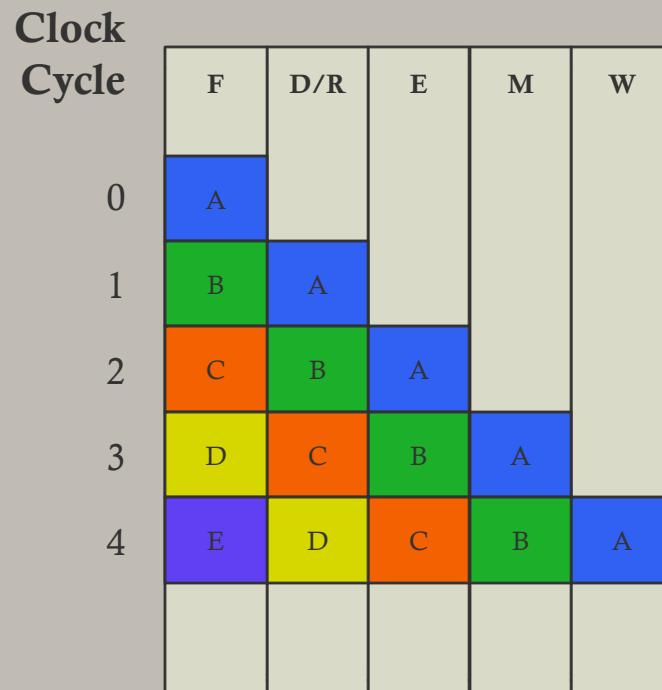
Pipelining

- In a non-pipelined CPU, instruction stages are idle much of the time.



Pipelining

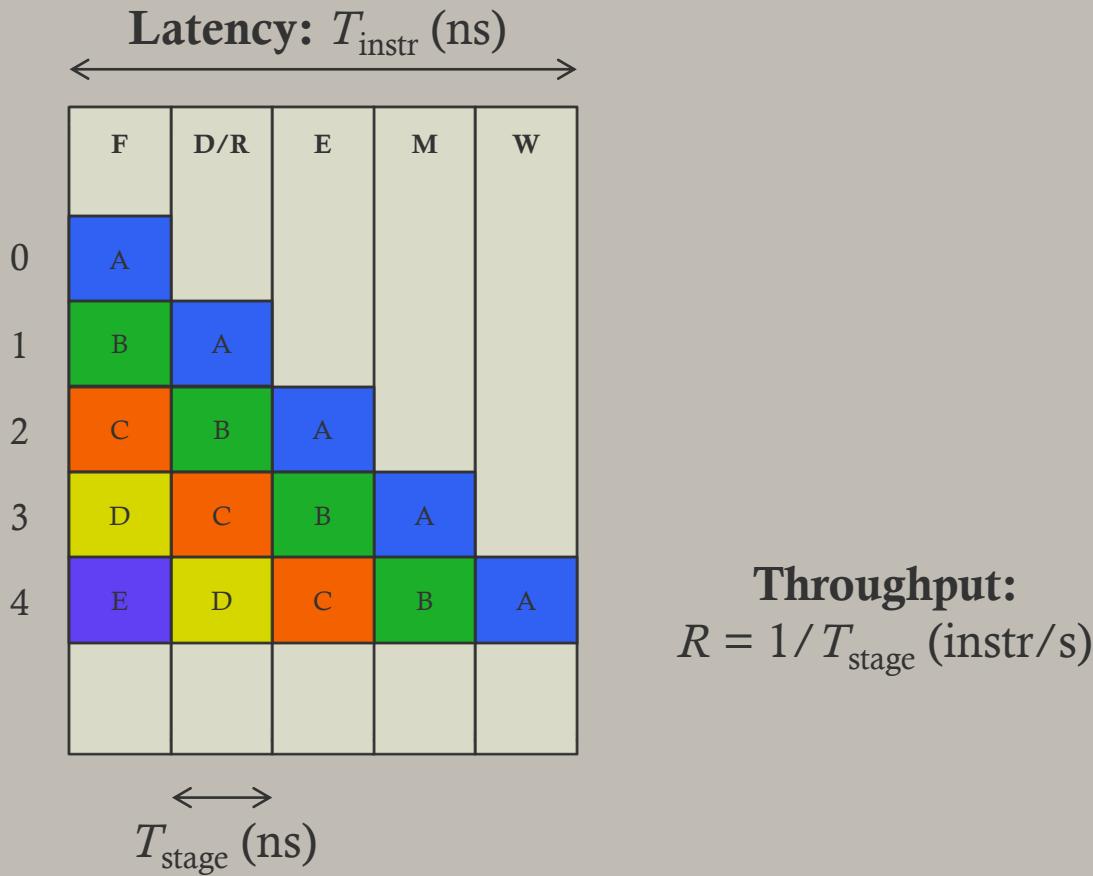
- To keep each stage busy, **issue** a new instruction every clock



Pipelining

- Pipelining is a form of **instruction-level parallelism** (ILP)
- Measuring performance
 - **Throughput:** Number of instructions retired per unit time
 - **Latency:** Length of time required to retire a single instruction

Pipelining



Pipelining

- A **stall** is defined as a situation in which not all stages are kept busy
 - What kinds of situations might cause stalls?

Discussion: 5 mins

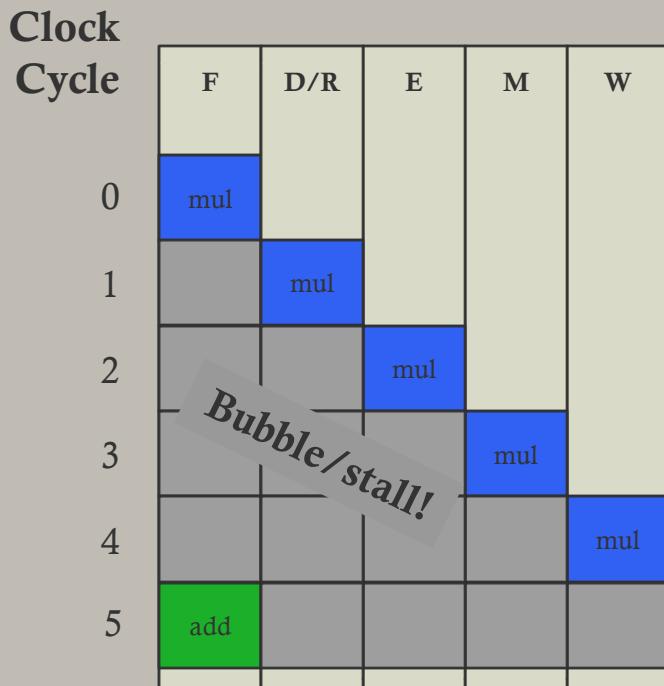
Pipelining

- A **stall** is defined as a situation in which not all stages are kept busy
 - In other words, the **next instruction** cannot be **issued** due to some kind of **dependency**
- Three kinds of dependencies in CPU pipeline:
 1. **Data** dependency: A register used in one instruction is used again in the next instruction
 2. **Branch** dependency: An instruction's operation is dependent on the results of the previous instruction (e.g. BZ)
 3. **Resource** dependency: Certain types of instructions can only execute on certain hardware (e.g. integer ALU)

Pipelining

- Dependencies are a property of an instruction stream
- Dependencies can result in **hazards**: Impacts on the CPU pipeline (stalls)
 1. **Data dependency** → **data hazard**: Results of a *previous instruction* are needed in order to issue this instruction
 2. **Branch dependency** → **control hazard**: Need to *predict branches* before result of conditional expression is known
 3. **Resource dependency** → **structural hazard**: Would like to issue, e.g., an integer add, but the ALU is *busy*

Data Hazards



```
int madd(int a, int b, int c)
{
    int temp = a * b;
    int result = temp + c;
    return result;
}

mul    r1, r2, r3 ; r1 = r2 * r3
add    r5, r1, r4 ; r5 = r1 + r4
```

Data Hazards

- Solution to data hazard problem:
 - Find something for the **CPU to do** during the bubble!
 - **Delay slots** are the instruction slots during which the CPU is idle waiting for the current instruction to complete
- Options:
 - **Programmer** can **manually reorder** statements in code
 - **Compiler** (optimizer) can reorder instructions automatically, based on analysis of program logic
 - **CPU** can reorder instructions on the fly
 - Called **out-of-order execution** (OOO execution)

Data Hazards

Clock Cycle	F	D/R	E	M	W
0	mul				
1	A	mul			
2	B	A	mul		
3	C	B	A	mul	
4	D	C	B	A	mul
5	add	D	C	B	A

OOO execution allows the CPU to **reorder** instructions A, B, C and D (taken from *before* the **add**) so that they come between the **add** and the **mul** instructions

```
mul    r1, r2, r3 ; r1 = r2 * r3
instrA r5
instrB r6, r7
instrC r0, r8
instrD r9
add    r5, r1, r4 ; r5 = r1 + r4
```

Data Hazards

<http://cseweb.ucsd.edu/classes/wi05/cse240a/pipe2.pdf>

- Data dependencies come in three flavors
 - **Read after write (RAW)**

mul **r1**, r2, r3
 ↓
add r5, **r1**, r4



- **Write after write (WAW)**

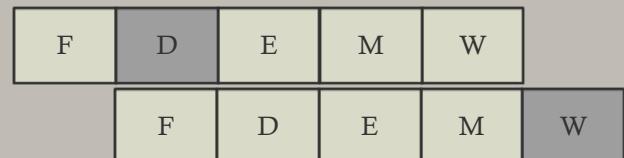
add **r1**, r2, r3
 ↓
sub **r1**, r2, r4



- **Write after read (WAR)**

- Can't happen in vanilla pipeline

mul r2, **r1**, r3
 ↓
sub **r1**, r5, r4



Data Hazards

- Exercise: Identify the RAW, WAW and WAR dependencies in the following instruction sequence

mul	r1,r2,r3
add	r5,r1,r4
mov	0x1E800,r5
mov	r6,0x1E750
mov	r6,(r6+4*r0)
mul	r1,r2,r0
mov	r2,#42
add	r1,r1,r2

Data Hazards

- Solution:

mul	$r1, r2, r3$	
add	$r5, r1, r4$	RAW
mov	$0x1E800, r5$	RAW
mov	$r6, 0x1E750$	
mov	$r6, (r6+4*r0)$	WAW, RAW
mul	$r1, r2, r0$	
mov	$r2, #42$	WAR
add	$r1, r1, r2$	

Data Hazards

- Exercise: Reorder the instruction sequence to **minimize** stalls

```
mul    r1,r2,r3
add    r5,r1,r4
mov    0x1E800,r5
mov    r6,0x1E750
mov    r6,(r6+4*r0)
mul    r1,r2,r0
mov    r2,#42
add    r1,r1,r2
```

Data Hazards

- One possible solution...
 - noticing that there are really two independent sequences
 - so interleave them

```
mul    r1,r2,r3  
add    r5,r1,r4  
mov    0x1E800,r5  
mov    r6,0x1E750  
mov    r6,(r6+4*r0)  
mul    r1,r2,r0  
mov    r2,#42  
add    r1,r1,r2
```



```
mul    r1,r2,r3  
mov    r6,0x1E750  
add    r5,r1,r4  
mul    r1,r2,r0  
mov    0x1E800,r5  
mov    r2,#42  
mov    r6,(r6+4*r0)  
add    r1,r1,r2
```

Register Renaming

- Register reuse within a sequence of instructions can introduce false dependencies (WAW, WAR)
- One solution: **register renaming**
 - Register names used by program (“virtual” registers) are **mapped** to real hardware registers by CPU or compiler

before register renaming...

add	r3 , r4, r5
ld	r7, (r3)
sub	r3 , r12, r11
st	(r15), r3

after register renaming...

add	hw3 , hw4, hw5
ld	hw7, (hw3)
sub	hw20 , hw12, hw11
st	(hw15), hw20

false dependency!

(WAR)

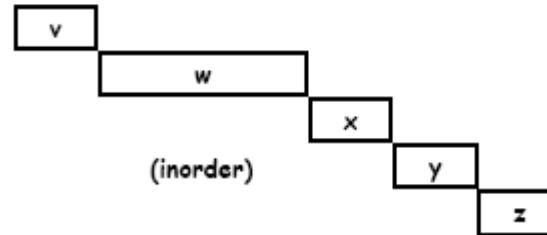


Register Renaming

http://www.math.nsysu.edu.tw/~lam/MPI/lecture/SUPERSCALAR_ARCHITECTURE.ppt

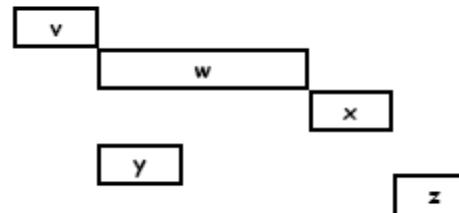
(Single adder, data dependence)

```
v:    addt $f2, $f4, $f10  
w:    mult $f10, $f6, $f10  
x:    addt $f10, $f8, $f12  
y:    addt $f4, $f6, $f4  
z:    addt $f4, $f8, $f10
```

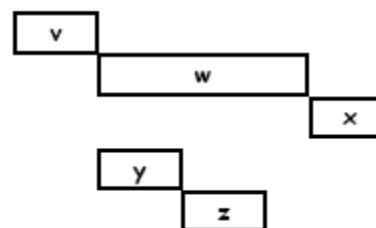


- Can start y as soon as adder available
- Must hold back z until \$f10 not busy & adder available

```
v:    addt $f2, $f4, $f10  
w:    mult $f10, $f6, $f10  
x:    addt $f10, $f8, $f12  
y:    addt $f4, $f6, $f4  
z:    addt $f4, $f8, $f10
```



```
v:    addt $f2, $f4, $f10a  
w:    mult $f10a, $f6, $f10a  
x:    addt $f10a, $f8, $f12  
y:    addt $f4, $f6, $f4  
z:    addt $f4, $f8, $f14
```



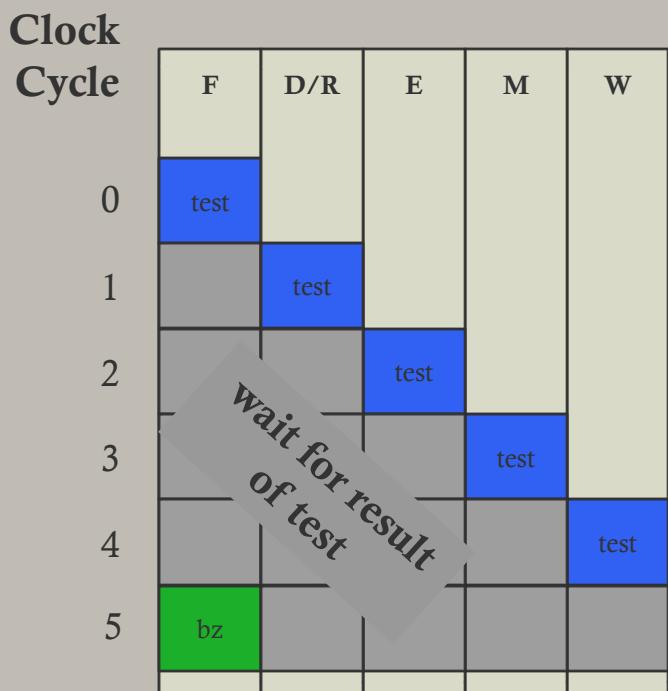
Memory Latency Hiding

- Note that stalls can also occur due to **memory latency**
 - `mov r3,[r1+8]`
 - Might take 4, 40 or 400 cycles! (L1, L2, main RAM)
- **OOO execution** can also help to **hide** memory latency
 - While waiting for memory controller to respond, execute other independent instructions in delay slots
- **Multithreading** is another way to **hide** memory latency
 - While one thread is waiting for memory, other thread(s) can be scheduled to use the CPU core

Control Hazards

- What should the CPU do when a **conditional branch** instruction is encountered?
 - Result of the condition *not yet known* at time of issue
 - Pipeline can be 10+ stages deep
 - Conditional branch therefore introduces 10+ cycle stall
- This is called a **branch penalty**
 - Control hazard
 - It's a special case of a RAW data hazard

Control Hazards



Control Hazards

- Solutions to branch penalty problem:
- Software:
 - **Loop unrolling** to increase run length (fewer if tests)
 - **Rearrange code** to minimize “branchiness”
- Hardware:
 - **Delay slots**
 - Choose instructions to run during the stall (**OOO execution**)
 - **Choose one** of the two branches and execute **speculatively**
 - If the guess is **correct**, no stall! ☺
 - If the guess is **wrong**, **flush** the pipeline, **back out** intermediate results, and **start again** on the correct branch

Speculative Execution

- How should the CPU guess?
 - Simple idea: Just **assume**...
 - *backwards* branches are *always* taken,
 - *forward* branches are *never* taken
 - Very common pattern in for loops

```
for (i = 0; i < count; ++i)
{
    // do work...

    if (special_case)
        break; // forward branch (rare)

} // backward branch (common)
```

Speculative Execution

- Another idea: Rely on manual tagging of **if** statements

```
for (i = 0; i < count; ++i)
{
    // do work...

    if (UNLIKELY(special_case))
        break;

}

// gcc definitions:
#define LIKELY(x)      __builtin_expect((x),1)
#define UNLIKELY(x)     __builtin_expect((x),0)
```

Speculative Execution

- Another idea: Include **branch prediction logic** on the CPU die
 - Keeps a history of which branches were taken
 - Logic can get quite complex
 - Some CPUs can detect **patterns** like Y, N, Y, N, Y, N, ...
 - Increases transistor count and circuit complexity significantly
- Speculative execution requires results of calculations to be stored temporarily in

Predication

- Yet another idea:
 - Don't branch at all!
 - Instead, always execute **both** sides of the branch
 - But **mark** each instruction with the condition (predicate) with which it is associated
 - CPU only **actually commits** results of predicated instructions when the condition (predicate) is true
- This approach is called **predication**
- Simpler form: partial predication
 - Conditional **move**, conditional **select** instructions

Predication

```
int SafeFloatDivide_pred(float a, float b, float d)

{
    // convert Boolean (b != 0.0f) into either 1U or 0U
    const unsigned condition = (unsigned)(b != 0.0f);

    // convert 1U -> 0xFFFFFFFFU
    // convert 0U -> 0x00000000U

    const unsigned mask = 0U - condition;

    // calculate quotient (will be QNaN if b == 0.0f)
    const float q = a / b;

    // select quotient when mask is all ones, or default value d when mask is all zeros (NOTE: this won't work as
    // written -- you'd need //to use a union to interpret the floats as unsigned for masking)

    const float result = (q & mask) | (d & ~mask);

    return result;
}
```

Partial Predication

- Conditional select
 - e.g., PowerPC **fsel** and **isel** instructions
 - isel** RR, RA, RB, *predicate_code*
 - If predicate bit is 1, move RA→RR
 - If predicate bit is 0, move RB→RR

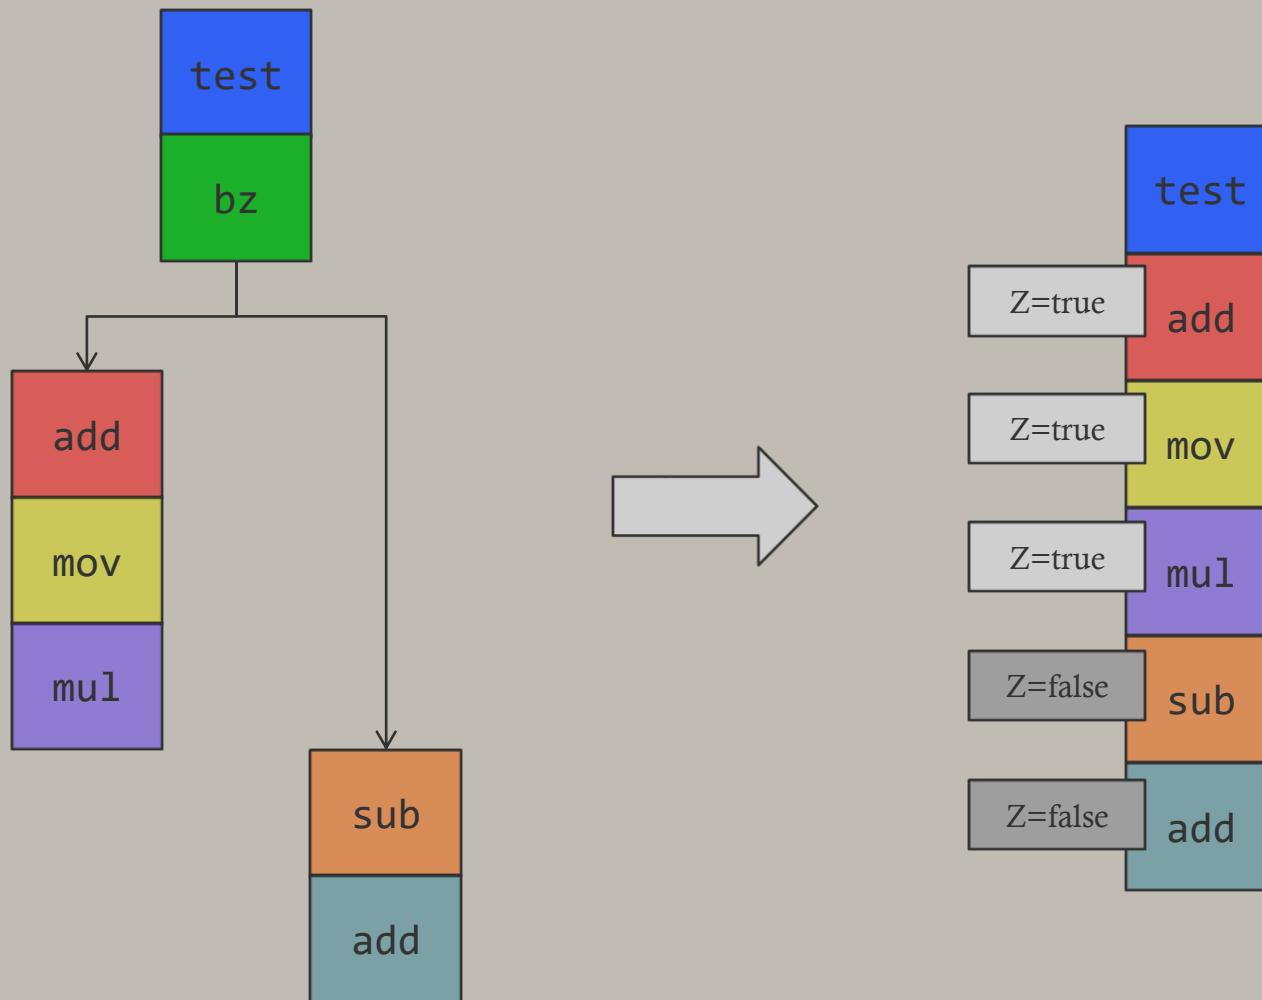
```
cmpwi %cr7,%r11,42
beq cr7,Label1 ; if r11 == 42
li %r0,0         ; else r0 = 0
b Label2
Label1:
li %r0,1         ; set r0 = 1
Label2:
```

```
li    %r1,0
li    %r2,1
cmpwi %cr7,%r11,42
isel %r0,%r2,%r1,30
```



30 means “check
the EQ status bit
of %cr7”

Full Predication



Superscalar

- If pipelining improves throughput, can we improve it further by adding **more parallelism**?
 - Yes: If CPU has **two of each** functional unit, it can issue **two instructions per clock**
- This is called **superscalar** CPU architecture

Superscalar

Superscalar

- Resource dependencies (structural hazards) can occur in a superscalar CPU
 - e.g., we'd like to issue 3 **integer arithmetic operations** consecutively, but there are only 2 integer ALUs
 - The third integer op must stall, or...
 - we must defer it and find a **non-integer op** to put in its place
 - floating-point arithmetic op
 - memory op
 - comparison op
 - SIMD vector op
 - etc.

Very Long Instruction Word (VLIW)

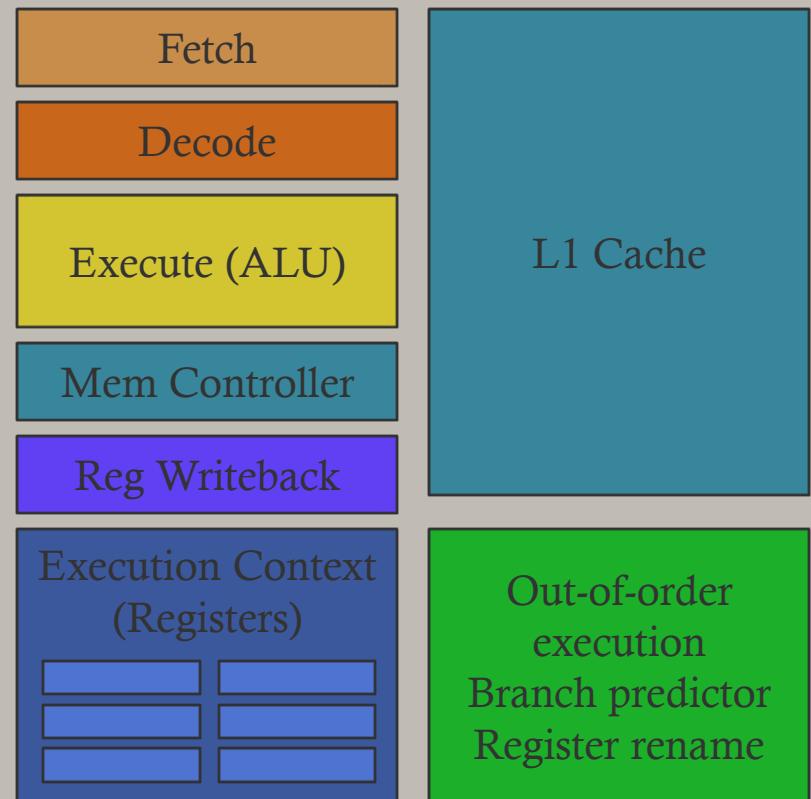
- Lots of silicon/transistors are required to implement:
 - OOO execution logic
 - Branch predictors
 - Register renaming
- Simpler design:
 - Include all the **parallel functional units**
 - ... but **omit** the complex **scheduling logic**
 - Leave it up to the **programmer/compiler** to find all opportunities for instruction-level parallelism...

Very Long Instruction Word (VLIW)

- But how to expose the parallel functional units to the programmer/compiler?
 - Each **instruction word** can encode **two or more instructions**
 - The program requests ILP **explicitly**
 - Gives rise to very long (wide) instruction words
 - **Very long instruction word** (VLIW) architecture
- e.g., VU0 and VU1 on the PlayStation2
 - Each VU had **two “slots”** per instruction
 - Up to the programmer/compiler to **reorder** assembly code in order to **fill both slots** of each instruction word

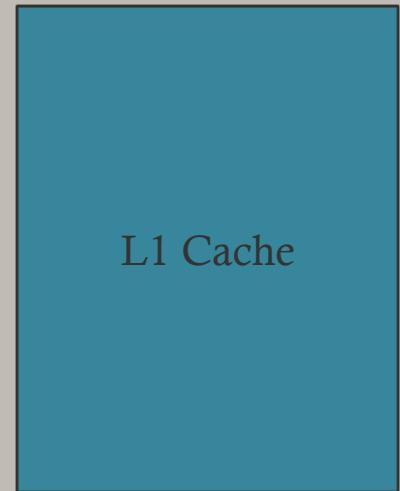
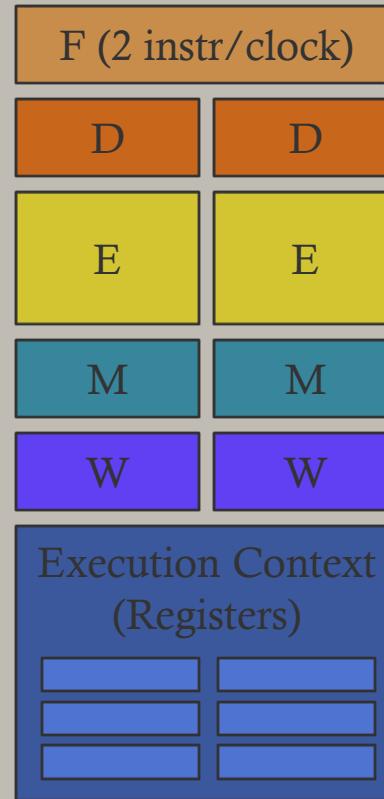
Pipelined CPU

opcode	AM	operand0	operand1
opcode	AM		
opcode	AM	operand0	
opcode	AM	operand0	operand1



Superscalar CPU

opcode	AM	operand0	operand1
opcode	AM		
opcode	AM	operand0	
opcode	AM	operand0	operand1

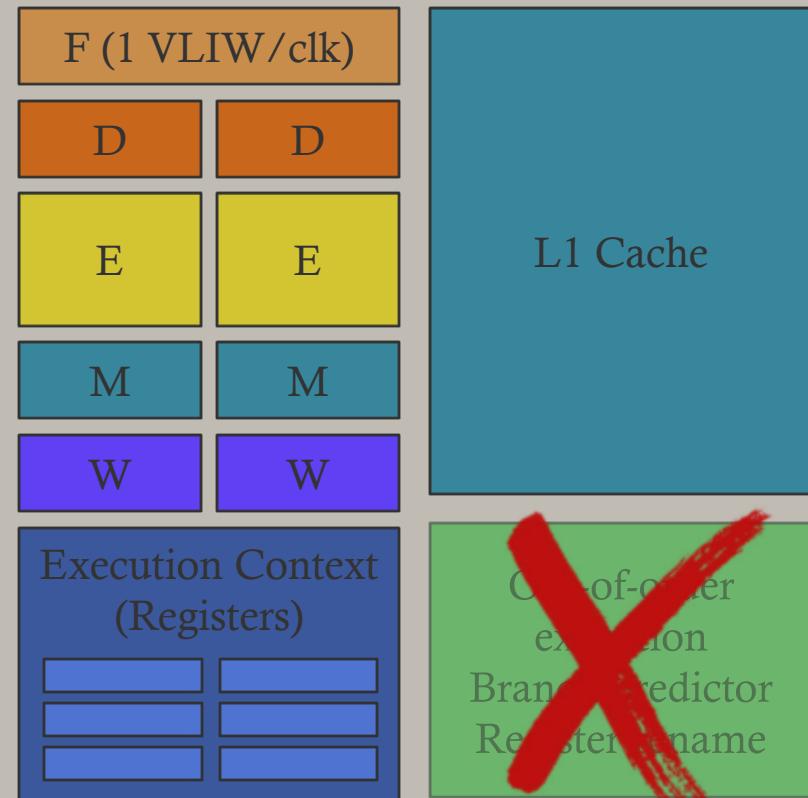


Out-of-order execution
Branch predictor
Register rename

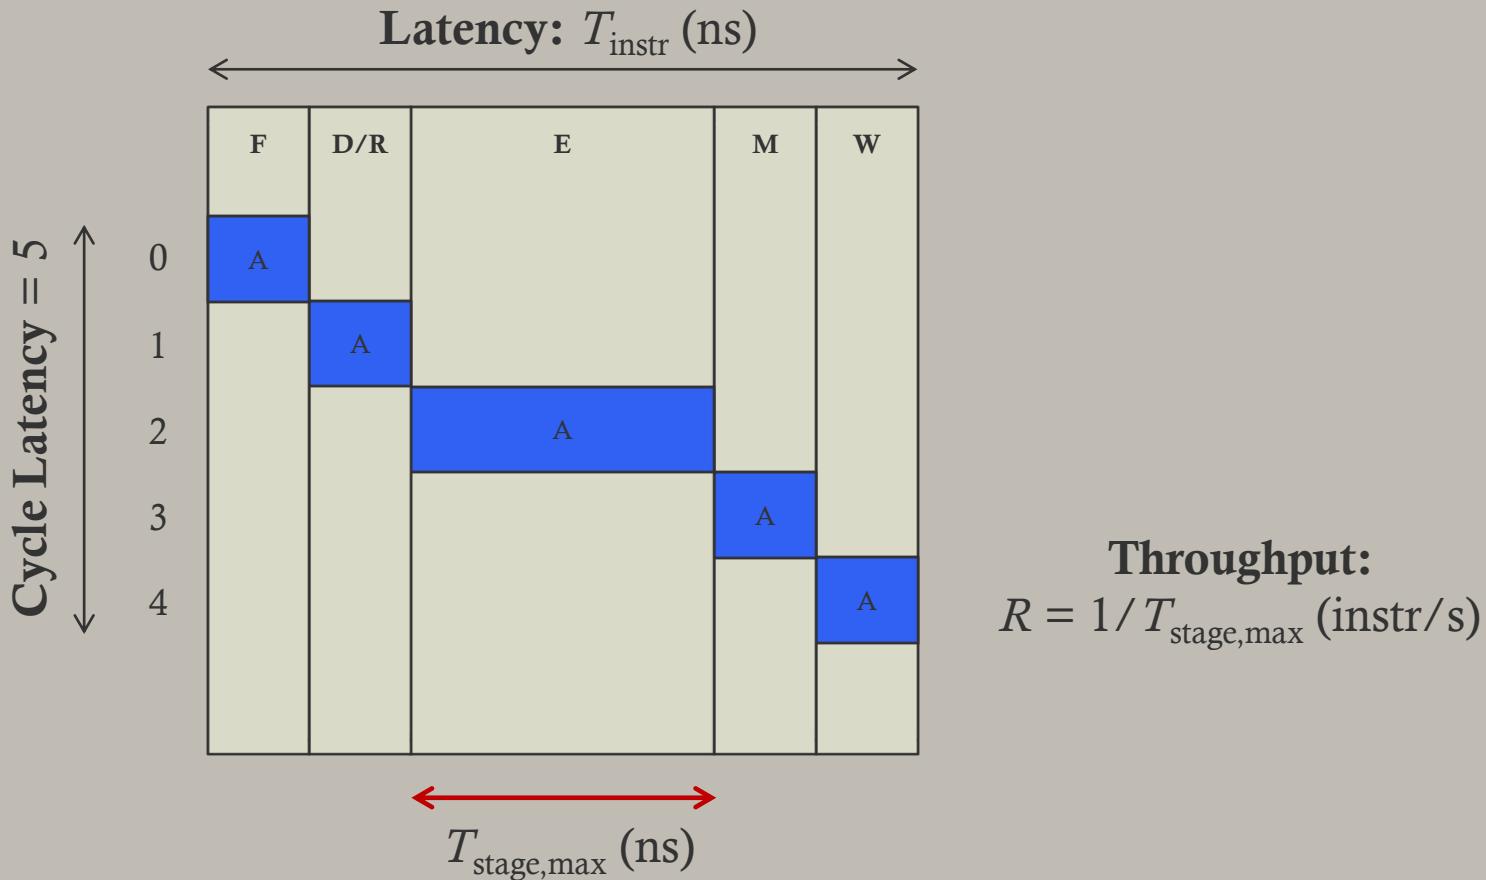
VLIW CPU

opcode	AM	operand0	operand1	opcode	AM	operand0	operand1
opcode	AM	operand0	operand1	opcode	AM	operand0	operand1

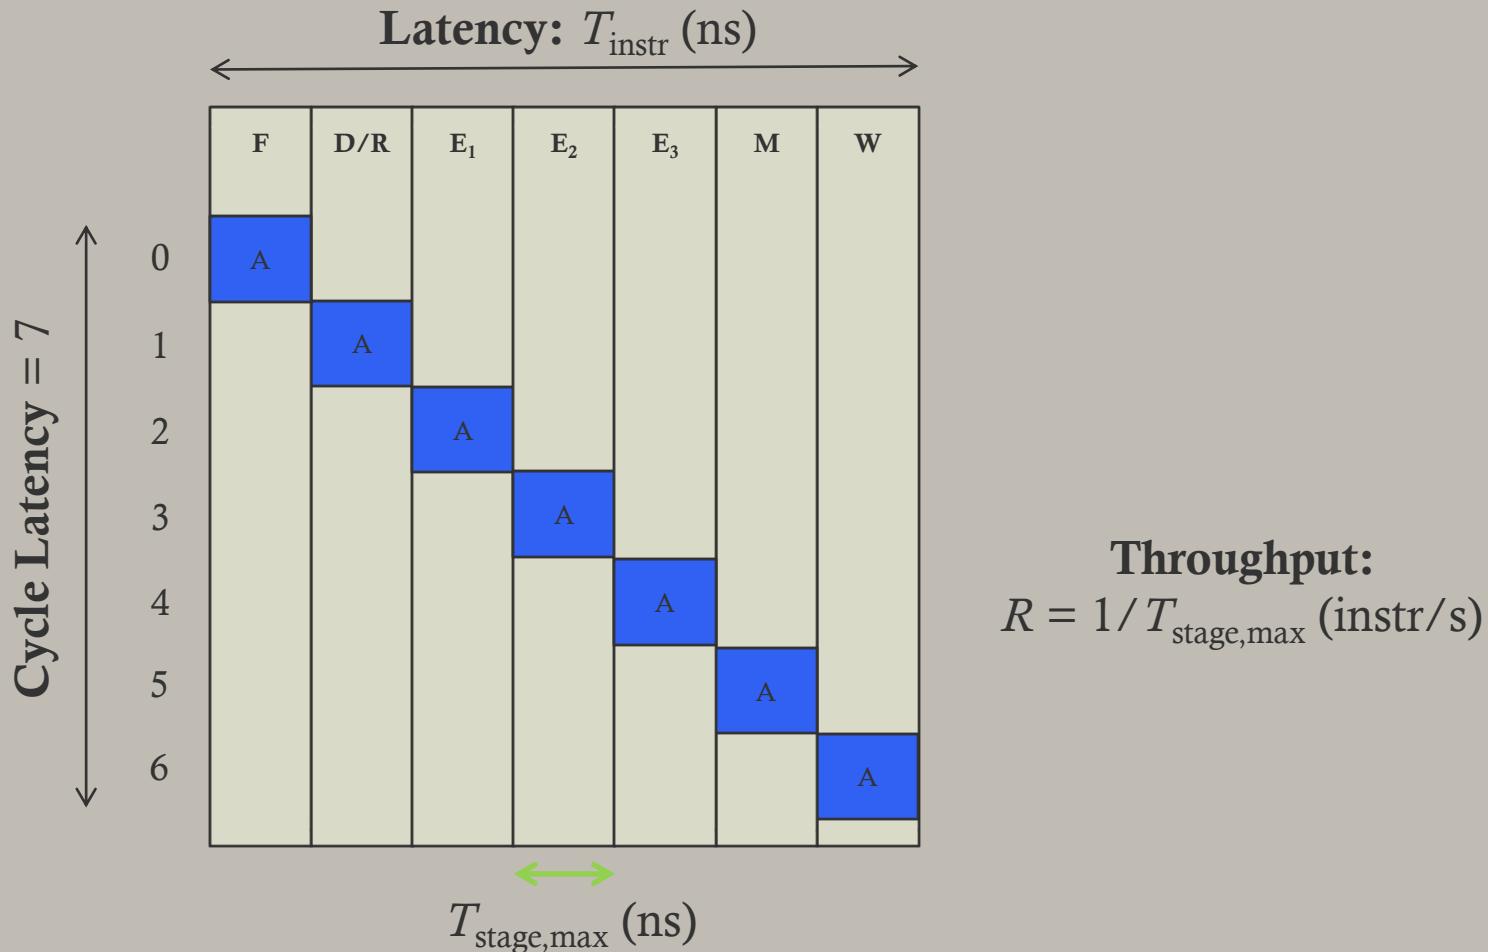
...



Toward Deeper Pipelines



Toward Deeper Pipelines

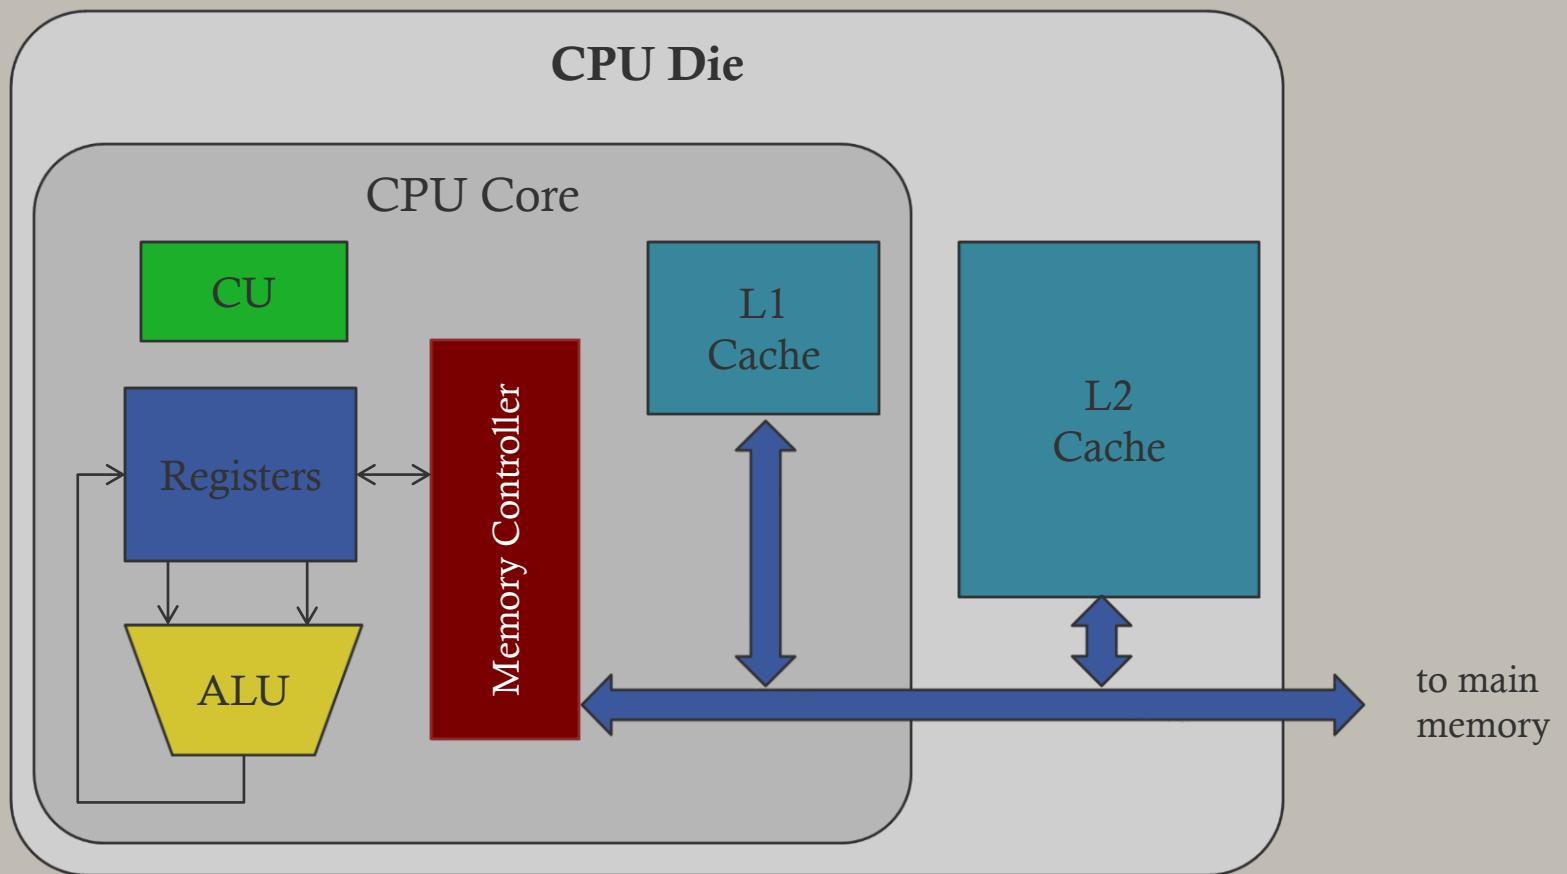


Toward Deeper Pipelines

- There are limits to how many stages a pipeline can be broken into
 - Increased clock speeds → increased heat and power consumption
 - More stages → more-complex forwarding logic
 - More-complex logic → more transistors, larger die area
 - More stages → bigger cost on branch mispredict
 - Minimum time to complete a stage is limited by technology considerations such as register access latency
- As such, pipeline depths *typically* vary between 5 stages at the low end, and 31 stages at the very high end

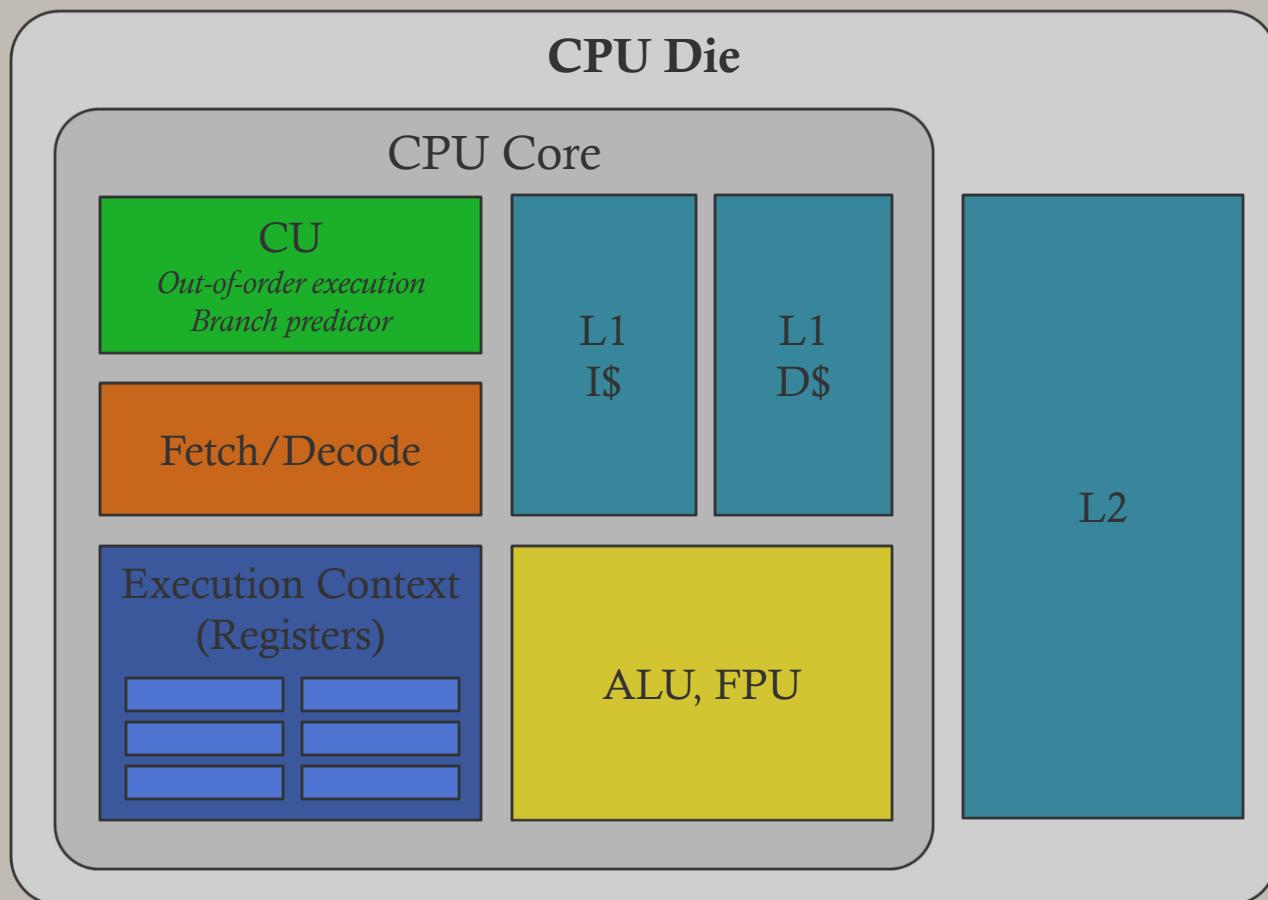
EXPLICIT PARALLELISM

Review: Simple Serial CPU

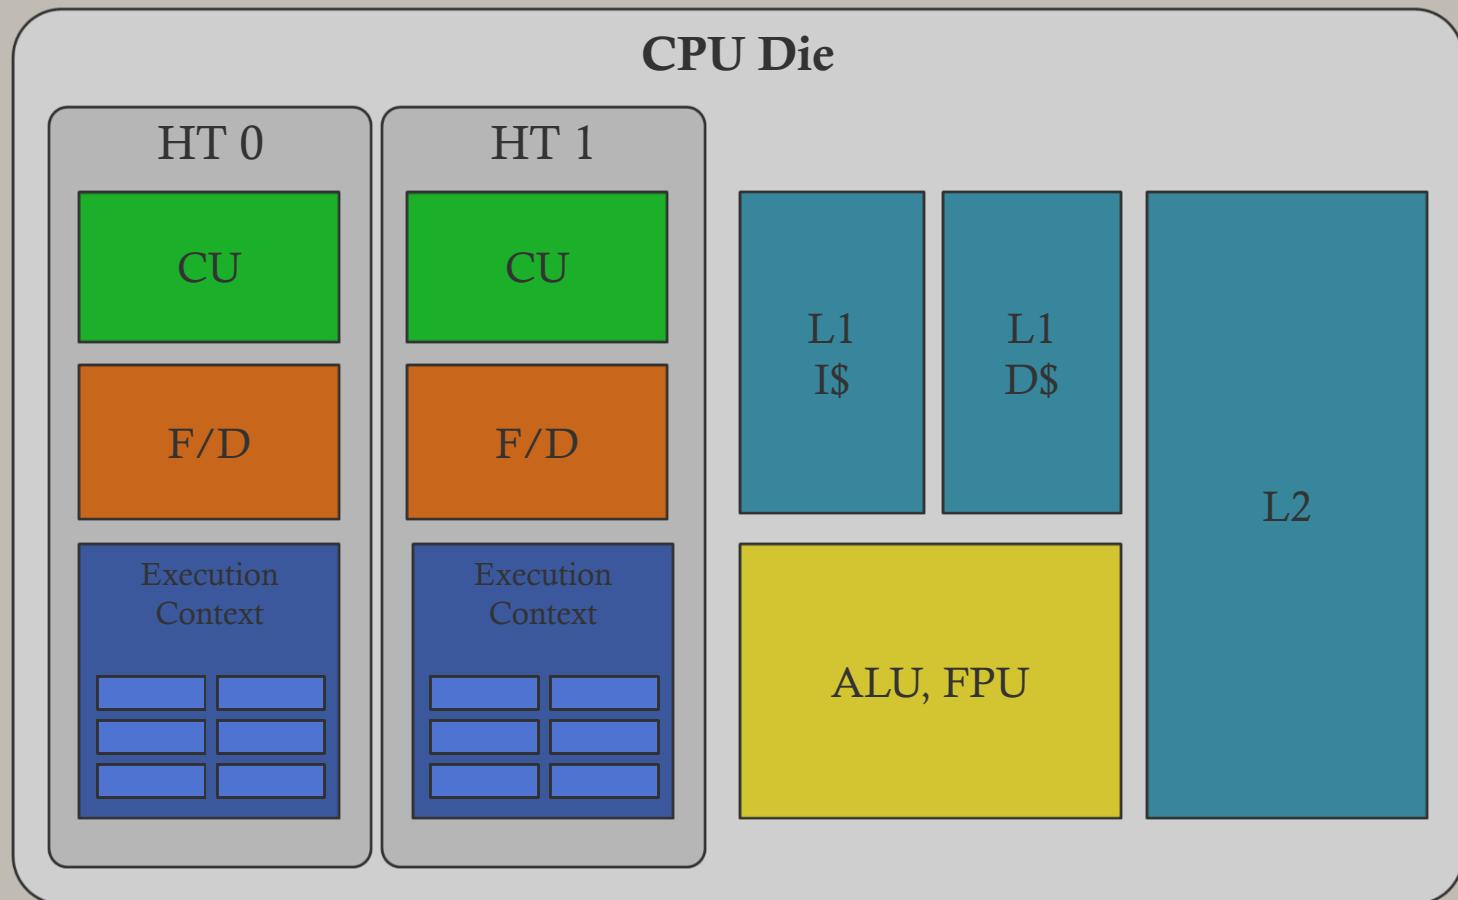


Review: Simple Serial CPU

- Rearrange the diagram a bit...



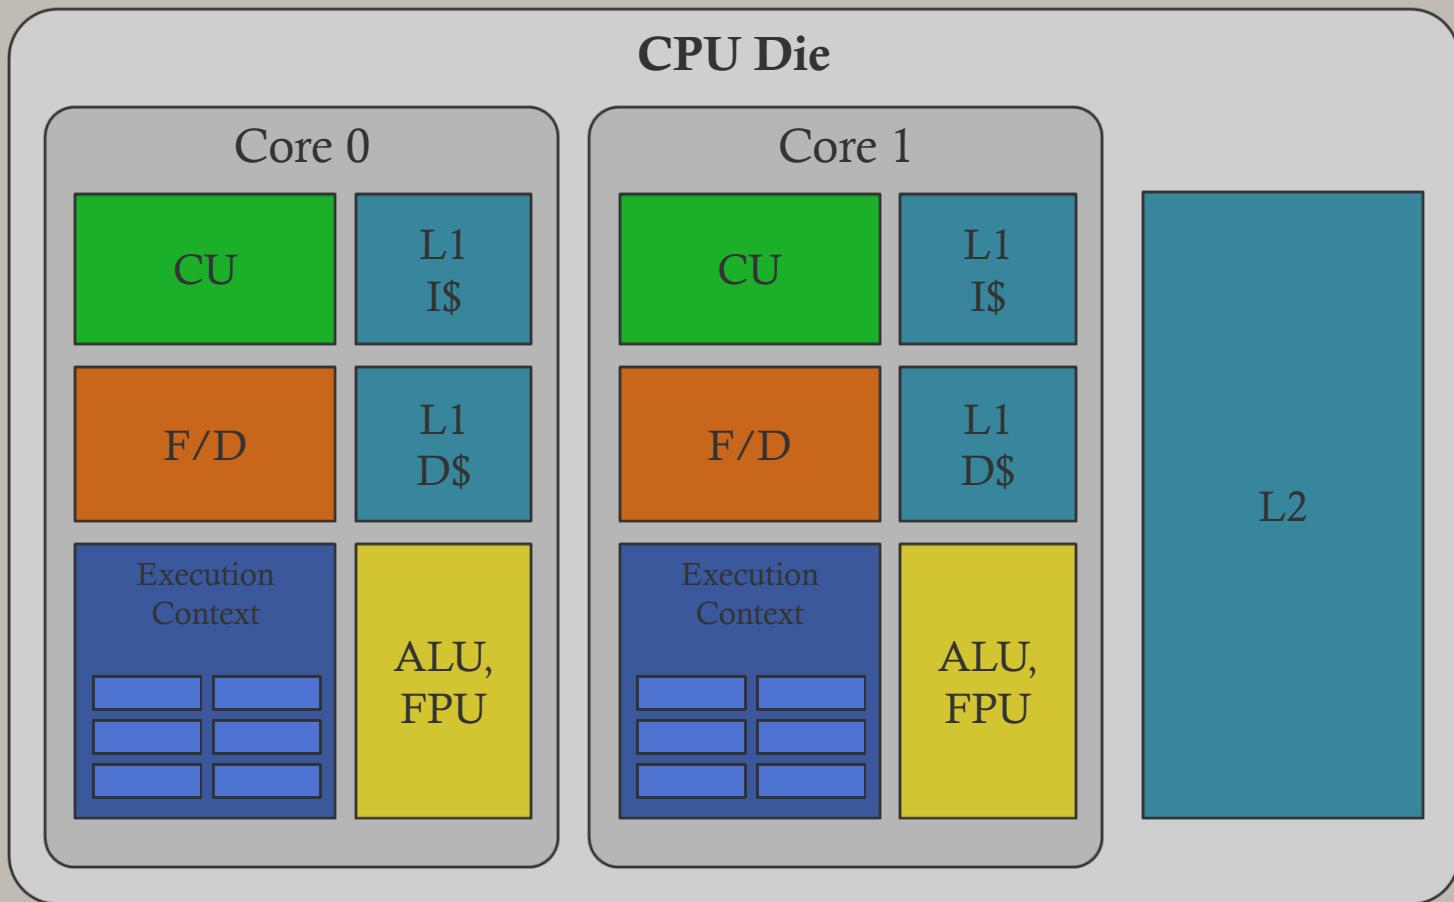
Hyperthreaded CPU



HyperThreading

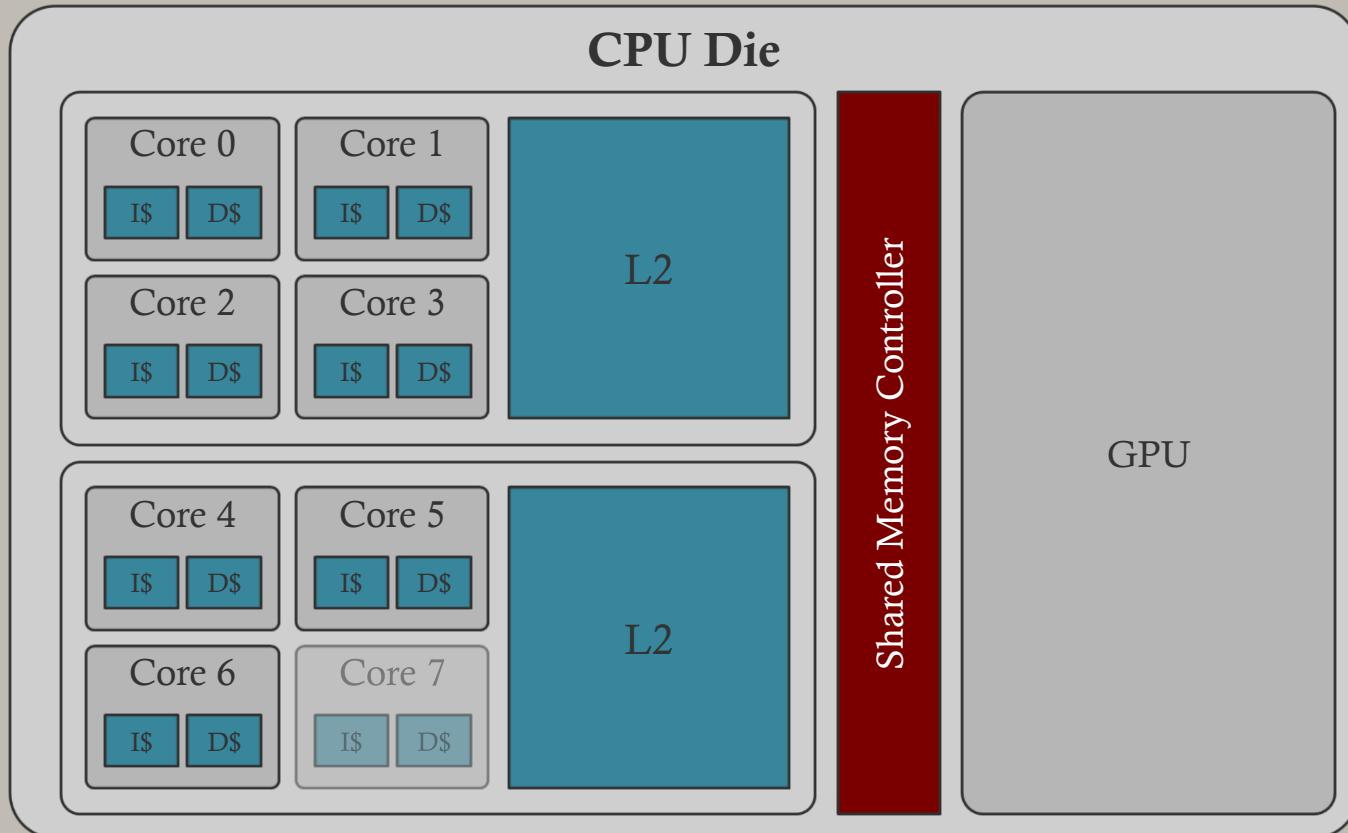
- If you have a windows OS, open task manager and go to CPU section, in that screen if you see the number of logical processors say 4 but your CPU is dual core, it is hyper threading.
- If you have Linux, open system monitor or whichever task manager you have on your Linux, and go to resources tab. Check how many CPUx do you see. If you see CPU1, CPU2 CPU3, CPU4 for a dual core CPU, it's hyperthreading.
- Run wmic (Windows Management Instrumentation) in command line
- CPU Get NumberOfCores,NumberOfLogicalProcessors /Format>List

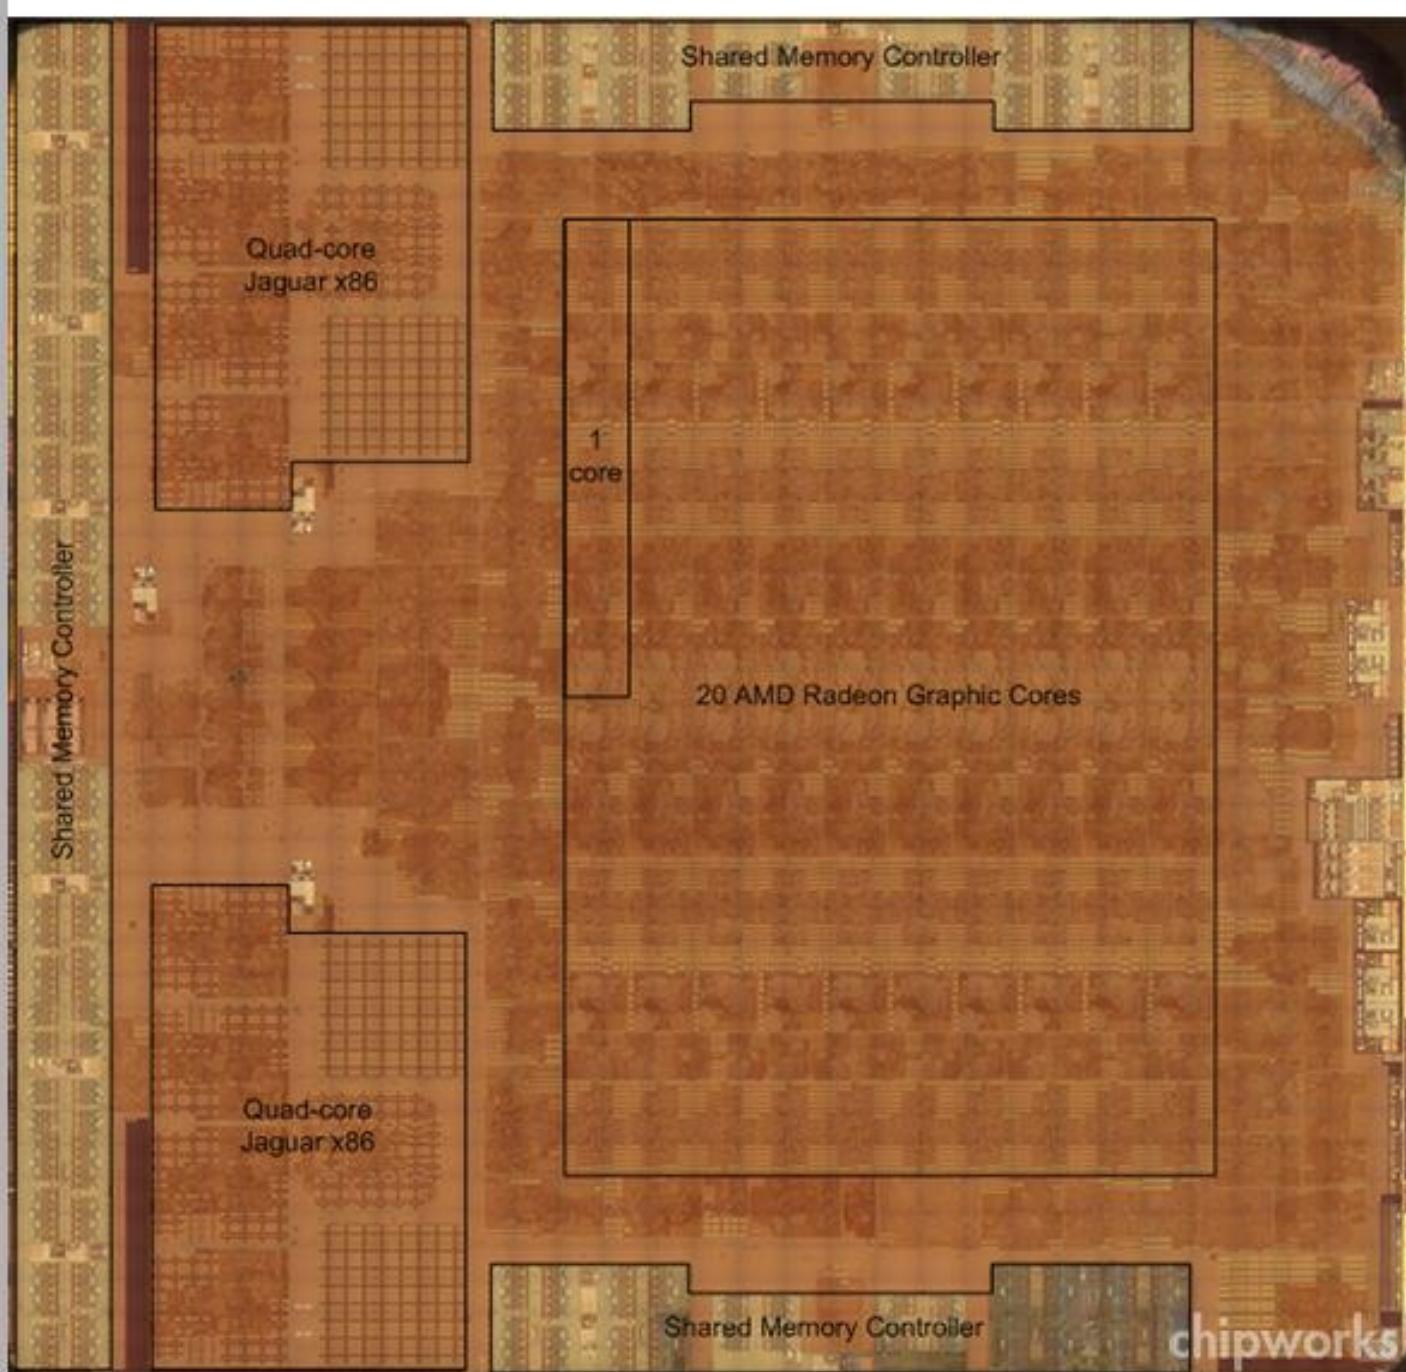
Multicore CPU



Multicore CPU

- e.g., PS4 architecture
 - 2 quad-core Jaguar CPUs
 - Radeon HD 7870(ish) GPU





Manycore GPU (SIMT)

- GPU architecture is a hybrid between SIMD and MIMD
- **SIMD cores** apply a single instruction stream (thread) to multiple data “lanes” in parallel
- To hide latency, many **threads** share a SIMD core via **time-slicing**
 - Nvidia calls this single instruction multiple **thread** (SIMT)
- We’ll dive into GPU architecture later (time permitting)...

Computer Clusters, Grids and the Cloud

- **Computer cluster:** Multiple homogenous computers, rack-mounted in close proximity, connected via fast local network
- **Grid computing:** Loosely coupled heterogeneous computers (e.g., SETI@home)
- **Cloud computing:** On-demand computing and storage
 - Software as a Service (SaaS)
 - *Insert-flavor-here* as a Service (XaaS)
 - Flexible, scalable, cost-effective
 - use only what you need, pay as you go